

Pointer Arithmetic and Lexical Scoping

CS449 Spring 2016

Review **Pitfall 1** from previous lecture

```
void foo(char *s) { s = "World"; }
```

```
int main()  
{  
    char *str = "Hello";  
    foo(str);  
    printf("%s\n", str);  
    return 0;  
}
```

- Problem here is that **s** becomes a copy of **str**.
- Both **s** and **str** contain location of “Hello”.
- After **foo** is finished, **s** contains location of “World”.
- **str** still has location of “Hello”.
- So, string “Hello” is **never** copied, only pointer to it is.

Review **Pitfall 1** from previous lecture

Solution:

```
void foo(char **s) { *s = "World"; }
```

```
int main()  
{  
    char *str = "Hello";  
    foo(&str);  
    printf("%s\n", str);  
    return 0;  
}
```

- We must send address of pointer **str**, which is copied to **s**.
- By dereferencing **s** we directly write to **str**.

Review **Pitfall 1** from previous lecture

Solution:

```
void foo(char **s) { *s = "World"; }
```

```
int main()
{
    char *str = "Hello";
    foo(&str);
    printf("%s\n", str);
    return 0;
}
```

- We must send address of pointer **str**, which is copied to **s**.
- By dereferencing **s** we directly write to **str**.

Pointers Review

- Pointer: Variable (or storage location) that stores the address of another location.
- Reference operator: e.g. `scanf("%d", &x);`
 - Address of “x” is passed in order to modify it
- Dereference operator: e.g. `*p = 0;`
 - Access the location pointed to by “p”
- Pointer to array vs. array of pointers
 - `char a[2][3]; char (*p)[3] = &a[1];`
 - vs. `char *p[2] = {"yes", "no"};`
- Value of array variable is the address to the first element of the array
 - Given `int a[3];`, `a == &a[0];` by definition
 - Thus can be stored in a pointer: `int *p = a; p[0] = 0;`
 - However, `a = p;` results in a compile error since `&a[0]` is not an l-value (a storage location)

Pointers Review (more)

Given array `int a[2][3][4]`, the following are valid

`int (*p)[3][4] = a;` (same as `int (*p)[3][4] = &a[0];`)

`int (*p)[4] = a[1];` (same as `int (*p)[4] = &a[1][0];`)

`int *p = a[1][0];` (same as `int *p = &a[1][0][0]`)

`int p = a[1][0][2];`

Derived types in C

- Types that are derived from fundamental data types.
 - Pointer
 - Array
 - Function
 - Structure
 - Union
- } We will cover these later

Size of Derived Types

```
#include <stdio.h>
#include <string.h>
int main()
{
    char buf[20];
    char *str = buf;
    strcpy(buf, "Hello");
    printf("str=%s, buf=%s\n", str, buf);
    printf("sizeof(str)=%u\n", sizeof(str));
    printf("sizeof(buf)=%u\n", sizeof(buf));
    printf("sizeof(\"Hello\")=%u\n", sizeof("Hello"));
    return 0;
}
```

```
>> ./a.out
str=Hello, buf=Hello
sizeof(str)=8
sizeof(buf)=20
sizeof("Hello")=6
```


Functions Review

- One way pointers are useful.. when passing pointer arguments to function
 - When modifying variables in caller's environment
 - E.g. `void swap(int *a, int *b);`
 - In a sense, allows the “return” of multiple values
 - When it's more efficient (to avoid copying)
 - E.g. `size_t strlen(const char* s);`
 - string “s” is not getting modified but pointer is passed instead of copying character array
 - What if we declare `size_t strlen(const char s[100]);`?
 - Will compile but same meaning as above. Don't bother using it. Arrays are always passed as pointers.

Const Qualifier

- What if you want to make sure a string passed as an argument is never modified?
- Declare string as const
 - e.g. `size_t strlen(const char* s);`
- Contract to caller that string pointed to by “s” will not be modified
- If contract is violated -> compile time error
 - With no contract -> runtime error (much harder to debug)
- Examples:
 - `const float pi = 3.14;`
 - `const char *str = “Hello”;` *//same as: `char const *str = “Hello”;`*
`const char *const str = “Hello”;`
 - Contract: String pointed to by “str” is immutable AND “str” cannot point to any other string

Example Use of Const

```
#include <stdio.h>
```

```
void foo(const char *s) { s[0] = 'h'; }
```

```
int main()
```

```
{
```

```
    const char *str = "Hello";
```

```
    foo(str);
```

```
    return 0;
```

```
}
```

Instead of error at runtime:

```
>> gcc ./main.c
```

```
>> ./a.out
```

```
Segmentation fault (core dumped)
```

You get error at compile time:

```
>> gcc ./main.c
```

```
./main.c: In function 'foo':
```

```
./main.c:3: error: assignment of  
read-only location '*s'
```

Example Use of Const

```
#include <stdio.h>
```

```
int main()
{
    const char *str = "Hello";
    char *str2 = str;
    str2[0] = 'h';
    return 0;
}
```

Compile time:

```
>> gcc ./main.c
./main.c: In function 'main':
./main.c:6: warning: initialization
discards qualifiers from pointer target
type
```

Run time:

```
>> ./a.out
Segmentation fault (core dumped)
```

- Assigning “const char*” to “char*” can potentially lead to violation of contract (Thus the warning.)

Pointer Arithmetic

- Another convenience offered by pointers
- Assuming “int a[3];”, following are equivalent:
 - a[2];
 - Get element 2 of int array “a”
 - *(a + 2);
 - Get element 2 offsets away from “a” or “&a[0]”
 - *(int*)((size_t)a + sizeof(int) * 2)
 - Direct address calculation of above
 - Instead of size_t you can use unsigned int for 32 bit systems and unsigned long for 64 bit systems.

Operations Permitted on Pointers

- Add constant offset (+, +=, ++)
 - E.g. “p = p + 1;”, “p += 1;”, “++p;”
- Subtract constant offset (-, -=, --)
- Offset between two pointers
 - E.g. “int offset = p1 - p2;”
- Comparison between two pointers. E.g.:
 - “p1 > p2” (p1’s offset is larger than p2’s offset)
 - “p1 == NULL” (if p1 is equal to NULL value)

Strcpy Using Pointer Arithmetic

```
char* strcpy(char *dest, const char *src) {  
    char *p = dest;  
    while(*p++ = *src++) ;  
    return dest;  
}
```

- Stops when *src == '\0' (when the null character at the end of src is reached)

Another Pointer Arithmetic Example

```
#include <stdio.h>
int main()
{
    int a[2][3];
    int *p = a[0];
    int (*p2)[3] = a;
    printf("p=%p, &a[0][0]=%p\n", p, &a[0][0]);
    printf("p2=%p, &a[0]=%p\n", p2, &a[0]);
    printf("p+1=%p, &a[0][1]=%p\n", p+1, &a[0][1]);
    printf("p2+1=%p, &a[1]=%p\n", p2+1, &a[1]);
    return 0;
}
```

```
>> ./a.out
p=0x7fff0c051ca0, &a[0][0]=0x7fff0c051ca0
p2=0x7fff0c051ca0, &a[0]=0x7fff0c051ca0
p+1=0x7fff0c051ca4, &a[0][1]=0x7fff0c051ca4
p2+1=0x7fff0c051cac, &a[1]=0x7fff0c051cac
```


Another Pointer Arithmetic Example

```
#include <stdio.h>
int main()
{
    int a[2][3];
    int *p = a[0];
    int (*p2)[3] = a;
    printf("p=%p, &a[0][0]=%p\n", p, &a[0][0]);
    printf("p2=%p, &a[0]=%p\n", p2, &a[0]);
    printf("p+1=%p, &a[0][1]=%p\n", p+1, &a[0][1]);
    printf("p2+1=%p, &a[1]=%p\n", p2+1, &a[1]);
    return 0;
}
```

- “p” and “p2” points to the same address
 - `p == &a[0][0]`
 - `p2 == &a[0]`
- “p+1” and “p2+1” point to different addresses
 - `p+1 == &a[0][1]` (base + `sizeof(int)`)
 - `p2+1 == &a[1]` (base + `sizeof(int[3])`)
- Why are pointers typed differently depending on base type?
 - To perform accurate pointer arithmetic
 - To know what code to generate on a dereference operation (what is the size of the data to be read?)
 - To enable type checking when an incompatible type is assigned

The void* Type

- Assigning to different pointer type results in compile error.
 - E.g. “int *p; char *p2 = p;” results in error
- Except when assigning to void* type
 - E.g. “int *p; void *p2 = p;” is perfectly fine
- Void pointer (void *)
 - Generic pointer representing any type
 - No casting needed when assigning to void* (vice versa)
 - Cannot be dereferenced / no pointer arithmetic
 - Size and type of variable pointed to not known
 - Used when the base type of a variable is unknown at compile time (typically will be known at runtime)

The NULL Value

- Equivalent to the numerical value “0”.
(Just like ‘\0’ is equivalent to “0”)
- NULL value means pointer points to nothing
- Make it a habit to initialize all invalid pointers to NULL. Advantages:
 - Can easily compare to NULL to check if pointer is valid
 - If accessing invalid pointer by mistake
 - Will result in a (clean) segmentation fault
 - Instead of accessing and corrupting some random memory

Command Line Arguments

```
#include <stdio.h>
int main (int argc, char *argv[])
{
    int i;
    printf ("command: %s.\n",argv[0]);
    for (i = 1; i < argc; i++) {
        printf("argv[%d] = %s\n", i, argv[i]);
    }
    return 0;
}
```

```
>> ./a.out foo bar
command: "./a.out".
argv[1] = foo
argv[2] = bar
```

- argc: total number of command line arguments (including command itself)
- argv: string array that contains the command line arguments

Lexical Scopes

- **Scope:** the portion of source code in which a symbol is legal and meaningful
 - **Symbol:** name of variable, constant, or function
 - At compile time, compiler matches each symbol to its corresponding memory location using scoping rules
- C defines four types of scopes
 - **Block scope:** within curly braces (e.g. within for loop)
 - **Function scope:** within functions
 - **Internal linkage scope:** within a single C source file
 - **External linkage scope:** global across entire program
- Means of **encapsulation** and **data-hiding**
 - In order to maximize encapsulation, minimize usage of globals

Lexical Scope Example

```
int global;  
static int file;  
int main()  
{  
    int function;  
    {  
        int block;  
    }  
}
```

<main.c>

```
extern int global;  
void foo() { global =  
10; }
```

<foo.c>

- “int block”: Block Scope
 - Only visible within curly braces
- “int function”: Function Scope
 - Only visible within “main()” function
- “static int file”: Internal Linkage Scope
 - Only visible within “main.c” file
 - **static**: storage class specifier limiting the scope
- “int global”: External Linkage Scope
 - Visible across entire program
 - In foo.c, declaring “extern int global” tells compiler “global” refers to a variable defined elsewhere
 - **extern**: storage class specifier declaring external linkage
 - Cannot declare a static global variable extern
- Also applies to functions except functions cannot be defined inside another function (only declared)

Shadowing

```
#include <stdio.h>
int n = 10;
void foo() {
    int n = 5;
    printf("Second: n=%d\n", n);
}
int main()
{
    printf("First: n=%d\n", n);
    foo();
    printf("Third: n=%d\n", n);
}
```

```
>> ./a.out
First: n=10
Second: n=5
Third: n=10
```

Shadowing

```
#include <stdio.h>
int n = 10;
void foo() {
    int n = 5;
    printf("Second: n=%d\n", n);
}
int main()
{
    printf("First: n=%d\n", n);
    foo();
    printf("Third: n=%d\n", n);
}
```

- Shadowing: when a variable in an inner scope “hides” a variable in an outer scope
- Function scope variable “int n = 5” shadows external linkage scope variable “int n = 10”
- Prevents local changes from inadvertently spilling over to global state
 - Whoever writes “foo()” does not need non-local knowledge of all the global variables in the program. This is important for modular programming.
- Do not over use shadowing
 - Reduces readability (have to think of scoping rules)
 - Mistakes while renaming variables may inadvertently expose a shadowed variable

Lifetime

- **Lifetime**: time from which a particular memory location is allocated until it is deallocated
 - Only applies to variables
 - Is a runtime property and describes behavior of program while executes (unlike scopes which is a compile time property)
- C defines three types of lifetimes
 - **Automatic**: automatically created and destroyed by code generated by compiler at scope begin and end
 - **Static**: allocated at program initialization and destroyed at program termination (static is an overloaded term)
 - **Manual**: manually created and destroyed by the programmer on the heap (will discuss this later)
- Allows efficient management of **memory** by compiler
 - Avoid static when possible to allow memory to be reclaimed
- Static variables are guaranteed to be **initialized to 0**
 - Done once by the Standard C Library at runtime before calling "main()"

Storage Classes

- **Storage class:** combination of variable scope and lifetime

Scope / Lifetime \ Block	Block	Function	Internal Linkage	External Linkage
Automatic	local	local	N/A	N/A
Static	static local	static local	static global	global

- **local:** Visible within curly braces and valid while executing code inside block or function
- **static global:** Visible within file and valid for entire duration
- **global:** Visible globally and valid for entire duration
- **static local:** (We haven't seen this yet) Visible within curly braces and valid for entire duration

Example of Wrong Storage Class

```
#include <stdio.h>

int* foo() {
    int x = 5;
    return &x;
}

void bar() { int y = 10; }

int main()
{
    int *p = foo();
    printf("*p=%d\n", *p);
    bar();
    printf("*p=%d\n", *p);
    return 0;
}
```

```
>> gcc ./main.c
./main.c: In function 'foo':
./main.c:4: warning: function
returns address of local variable
>> ./a.out
*p=5
*p=10
```

Example of Wrong Storage Class

```
#include <stdio.h>

int* foo() {
    int x = 5;
    return &x;
}

void bar() { int y = 10; }

int main()
{
    int *p = foo();
    printf("*p=%d\n", *p);
    bar();
    printf("*p=%d\n", *p);
    return 0;
}
```

- What happened?
- The lifetime of “int x” is within function foo()
- When foo returns, it returns a pointer to a deallocated memory location
- When *p is printed for the first time, it accesses the deallocated location but the location has not been overwritten yet, so it prints the correct value 5
- When function bar() is called, variable “int y” reuses the same location as “int x” that has been deallocated, overwriting the value with 10
- When *p is printed for the first time, it prints the overwritten value 10
- Why “int x” and “int y” end up in the same location will become clearer when we later talk about activation records and function stacks

Fix Using Static Local Storage Class

```
#include <stdio.h>
int* foo() {
    static int x = 5;
    return &x;
}
void bar() { int y = 10; }
int main()
{
    int *p = foo();
    printf("*p=%d\n", *p);
    bar();
    printf("*p=%d\n", *p);
    return 0;
}
```

```
>> gcc ./main.c
>> ./a.out
*p=5
*p=5
```

Returning String Using Static Local

```
char *asctime(const struct tm *timeptr) {  
    static char wday_name [7][3] = { "Sun", "Mon", "Tue", "Wed",  
        "Thu", "Fri", "Sat" };  
  
    static char mon_name [12][3] = { "Jan", "Feb", "Mar", "Apr",  
        "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov", "Dec" };  
  
    static char result [26];  
  
    sprintf(result , "%.3s %.3s%3d %.2d:%.2d:%.2d %d\n",  
        wday_name[timeptr->tm_wday], mon_name[timeptr->tm_mon],  
        timeptr->tm_mday, timeptr->tm_hour, timeptr->tm_min,  
        timeptr->tm_sec , 1900 + timeptr->tm_year);  
  
    return result;  
}
```

- “result” string still valid even after function returns

Keeping Track of Internal State Using Static Local

```
void foo() {  
    static unsigned int count = 0;  
    printf("Foo called %d time(s).\n", ++count);  
}
```

- “count” is a property of function foo(), and as such is properly encapsulated through scoping
- But persistence of “count” allows tracking of internal state across calls to foo

Keeping Track of Internal State Using Static Local

```
int main() {  
    char str[] = "Blue,White,Red";  
    char *tok = strtok(str, ",");  
    while(tok != NULL) {  
        printf("token: %s\n", tok);  
        tok = strtok(NULL, ",");  
    }  
    return 0;  
}
```

```
>> ./a.out  
token: Blue  
token: Red  
token: White
```

- char *strtok(char *str, const char *delim): parses “str” into a sequence of tokens using “delim” as delimiter (returns a token at each call)
- First call to strtok: return first token by replacing first occurrence of “delim” in “str” with ‘\0’ null character
- Subsequent calls to strtok (with NULL argument): return next token by picking up where we left off in “str” and replacing next occurrence of “delim” with ‘\0’ null character
- Location of next token is kept across calls using a static local variable

Register Storage Class Specifier

- Third storage class specifier (besides “static” and “extern”)
 - e.g. “register int x;”
- Tells the compiler to allocate the variable in a CPU register rather than memory, if possible
- Speeds up access for frequently used variables
- Can only be used for local (automatic) variables, which have a finite lifetime
 - For static variables, cannot allocate location in a register for entire duration of program
- Mostly obsolete and ignored by compiler since register allocators do a better job nowadays

Volatile Storage Class Specifier

- Fourth storage class specifier (besides “static”, “extern”, and “register”)
 - e.g. “volatile int x;”
- Tells the compiler to never allocate the variable in a register and always in memory
- Used for hardware devices when memory mapped I/O is required
- (Often **mistakenly**) used for multi-threaded programming to force communication through memory
 - Not recommended since volatile does not establish a “*happens-before*” relationship. We will talk more about “*happens-before*” in the synchronization section.