

Memory Management

CS449 Spring 2016

Lifetimes

- Lifetime: time from which a particular memory location is allocated until it is deallocated
- Three types of lifetimes
 - Automatic (within a scope)
 - Static (duration of program)
 - Manual (explicitly controlled by programmer)
- Manual control involves inserting into your code
 - Allocation calls
 - Deallocation calls

Manual Allocation Pros/Cons

- Pros (Flexibility)
 - Can create storage locations on demand without declaring them as variables
 - Useful for constructing dynamic data structures (e.g. linked lists, trees, graphs) whose size and shape can change depending on user input
- Cons (Complexity)
 - Programmer has to think about the behavior of program to determine the lifetimes of locations

C Standard Library Functions

- Takes care of the nitty-gritty details of memory management
 - System calls to request more memory from OS
 - Keeping track of areas of allocated memory
 - Keeping track of areas of free memory
 - Searching for suitable area to allocate memory
- Programmer only has to worry about *when* to allocate/deallocate and not *how*
- Just include functions declared in `<stdlib.h>`

Allocation Functions

- `void *malloc(size_t size)`
 - Allocates `size` bytes and returns a pointer to the allocated memory (NULL on failure)
 - Memory is not cleared
- `void *calloc(size_t nmemb, size_t size)`
 - Allocates `nmemb * size` bytes and returns a pointer to the allocated memory (NULL on failure)
 - Memory is set to 0
- Check return value for out-of-memory error

Deallocation Function

- `void free(void *ptr)`
 - Frees the memory space pointed to by `ptr`, which must have been previously allocated
 - Counterpart for both `malloc()` and `calloc()`
 - If `free()` has already been called on the same location, behavior is undefined

Re-Allocation Function

- `void *realloc(void *ptr, size_t size)`
 - Changes size of memory block pointed to by `ptr` to `size` bytes and returns that pointer (NULL on failure)
 - Newly allocated memory will be uninitialized
 - If `ptr` is NULL, same as `malloc`
 - If `size` is zero, and `ptr` is not NULL, same as `free`
- Use when you want to resize previously allocated memory without losing its contents

Malloc/Free Example

```
#include <stdio.h>
#include <stdlib.h>
```

```
int main(int argc, char *argv[])
{
    int *p, i, length = atoi(argv[1]);
    p = (int *)malloc(length * sizeof(int));
    srand((unsigned int)time(NULL));
    for(i = 0; i < length; i++)
        p[i] = rand();
    for(i = 0; i < length; i++)
        printf("%d \n", p[i]);
    free(p);
    return 0;
}
```

```
>> ./a.out 5
729751416
264693780
884704288
90101471
690008936
```


Malloc/Free Example

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    int *p, i, length = atoi(argv[1]);
    p = (int *)malloc(length * sizeof(int));
    srand((unsigned int)time(NULL));
    for(i = 0; i < length; i++)
        p[i] = rand();
    for(i = 0; i < length; i++)
        printf("%d \n", p[i]);
    free(p);
    return 0;
}
```

- Calculate the number of bytes of memory you want to allocate using `sizeof`
- Don't forget to free memory once you were done using it
- The lifetime of the locations allocated using `malloc()` is independent of pointer `p`
 - Pointer `p` is deallocated at the end of scope `main()`
 - The malloced locations lifetime ends on `free()`

Memory Management Strategy

- Allocation is relatively easy (just like `new` in Java): allocate memory before using it
- Basic Rules for deallocation
 - Memory should be freed when it can no longer be reached by the program
 - That is, when the *last pointer* referencing it is updated
 - Freeing before: can lead to accessing freed memory
 - Freeing after: last reference is gone, so no way to free!
 - This is actually the rule used by the Java garbage collector
 - How do you know if you are the last pointer?
 - Reference counting (how many point to you) can be helpful

Common Memory Errors

- Memory Leak
 - Forgetting to free unused memory after pointer update
 - Result: Steady rise in memory consumption leading to degraded performance and eventual out-of-memory error
- Double Free
 - Freeing the same memory location twice
 - Result: Undefined. Depends on C stdlib implementation.
- Dangling Pointer
 - Accessing memory that has already been freed
 - Result: Potential memory corruption when that memory is allocated for something else
- Out-of-bounds Access
 - Accessing memory beyond the allocation boundary
 - Result: Potential memory corruption when memory beyond boundary is allocated for something else
- Memory errors are typically Heisenbugs (non-deterministic bugs)

Valgrind

- Diagnostic tool that detects common memory errors (among other things) at runtime
- Command: `valgrind [options] <program>`
- Not perfect.
 - Can miss errors (sometimes)
 - Can report errors when there are none (rarely)
- Not a replacement for sound programming

```
#include <stdio.h>
#include <stdlib.h>
int main ()
{
    char *p = malloc(100);
    return 0;
}
```

Valgrind Example

```
>> gcc -g main.c
>> ./a.out
>> valgrind --leak-check=full --track-origins=yes ./a.out
==32563== HEAP SUMMARY:
==32563==    in use at exit: 100 bytes in 1 blocks
==32563== total heap usage: 1 allocs, 0 frees, 100 bytes allocated
==32563==
==32563== 100 bytes in 1 blocks are definitely lost in loss record 1 of 1
==32563==    at 0x4A069EE: malloc (vg_replace_malloc.c:270)
==32563==    by 0x4004DC: main (main.c:5)
```

- The “-g” option given to GCC inserts debug symbols to binary, enabling valgrind to locate the errors more accurately

How does malloc interact with the OS?

- C library File I/O functions were built on top of 5 system calls: open, read, write, lseek, close
- C library memory management functions are built on top of 2 system calls: brk and mmap
- `int brk(void *addr)`
 - Changes program break (location of the end of the process's data segment)
 - Increasing program break → allocates memory
 - Decreasing program break → deallocates memory

Tracing System Calls for malloc

- `strace ./a.out`
(a.out is the text file dumper run on a Hello World main.c file)

```
(116) thot $ strace ./a.out 2> /tmp/a.err && cat /tmp/a.err
open("main.c", O_RDONLY)          = 3
read(3, "#include <stdio.h>\nint main()\n{\n"... , 4096) = 75
write(1, "#include <stdio.h>\n", 19)  = 19
write(1, "int main()\n", 11)         = 11
write(1, "{\n", 2)                  = 2
write(1, " printf(\"Hello world!\\n\");\n", 28) = 28
write(1, " return 0;\n", 12)        = 12
write(1, "}\n", 2)                  = 2
write(1, "\n", 1)                   = 1
read(3, "", 4096)                   = 0
close(3)                            = 0
```

- Notice the difference in buffering for “read” and “write”

Tracing System Calls for malloc

```
#include <stdlib.h>
```

```
int main(int argc, char *argv[])  
{  
    int i = 0;  
    for(i = 0; i < 100; i++) {  
        malloc(4096);  
    }  
    return 0;  
}
```

```
>> strace ./a.out
```

```
[sic]
```

```
brk(0)                                =  
0x601000
```

```
brk(0x623000)                        =  
0x623000
```

```
brk(0x644000)                        =  
0x644000
```

```
brk(0x665000)                        =  
0x665000
```

```
brk(0x686000)                        =  
0x686000
```

```
[sic]
```


Practice Problem 1

- Write a program in C, which reads from standard input an integer ***N*** and generates an array of ***N*** random numbers;
- Write a function in C, which gets an array of random numbers as input parameter and returns the number of occurrences of the second biggest/smallest number in the array.
- Example: ***{-1, 0, 1, 2, 0, 0, 1, 1, 1, 1, -1}***; ***3*** for the second smallest and ***5*** for the second biggest.

Practice Problem 2

- Write a library in C, which implements rational numbers. The numbers should be represented as a ratio of numerator (integer) and denominator (integer). Write functions for the following operations of two rational numbers – addition, subtraction, multiplication, division, comparison, simplification.

Practice Problem 3

- Write a program in C, which reads from standard input an integer **N** , generates an 2 arrays of **N** integer numbers – **A_1, A_2, \dots, A_N** and **B_1, B_2, \dots, B_N** . The numbers in the 2 arrays should be read from the standard input. They are the catheti of **N** right-angled triangles.
- Write a function in C, which finds the hypotenuse of a triangle. Use this function to find the index of a triangle with the longest hypotenuse.
- You can use function: **`double sqrt(double x)`** from **`math.h`** library to calculate square root.

Practice Problem 4

- Write a function in C, which gets an integer ***N*** as input parameter and returns an array of 3-digit numbers for which the sum of the digits is equal to ***N***.
- Example: ***N=2; {110, 200}***
- *This problem can be solved using recursion*

Practice Problem 5

- Write a function in C, which gets an integer ***N*** as input parameter. The function should read from standard input numbers until their sum exceeds ***N***. Once ***N*** is exceeded, the function should print the numbers in sorted order.
- N.B. Negative numbers are valid input.