

Executables and Linking

CS449 Spring 2016

Remember External Linkage Scope?

```
#include <stdio.h>
int global = 0;
void foo();
int main()
{
    foo();
    printf("global=%d\n", global);
    return 0;
}

<main.c>
```

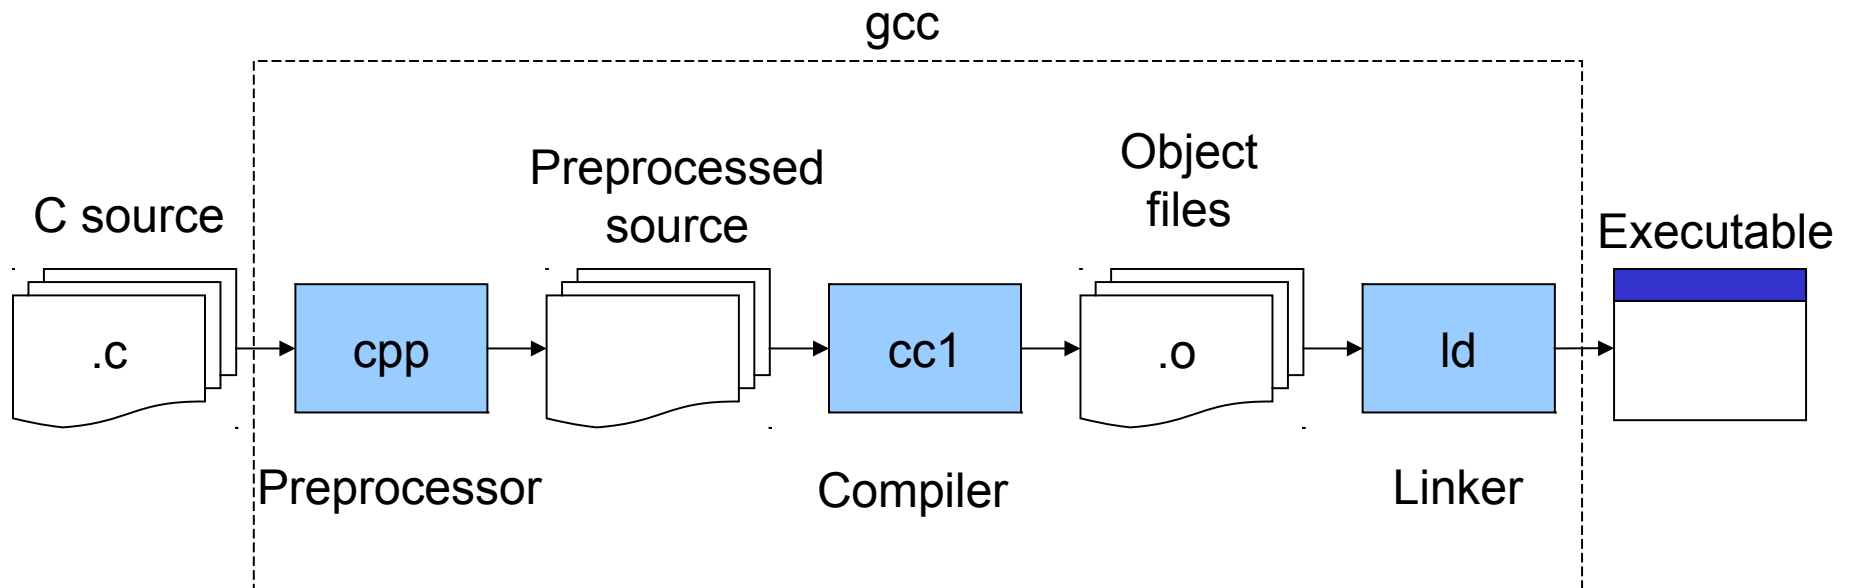
```
extern int global;
void foo() { global += 10; }

<foo.c>
```

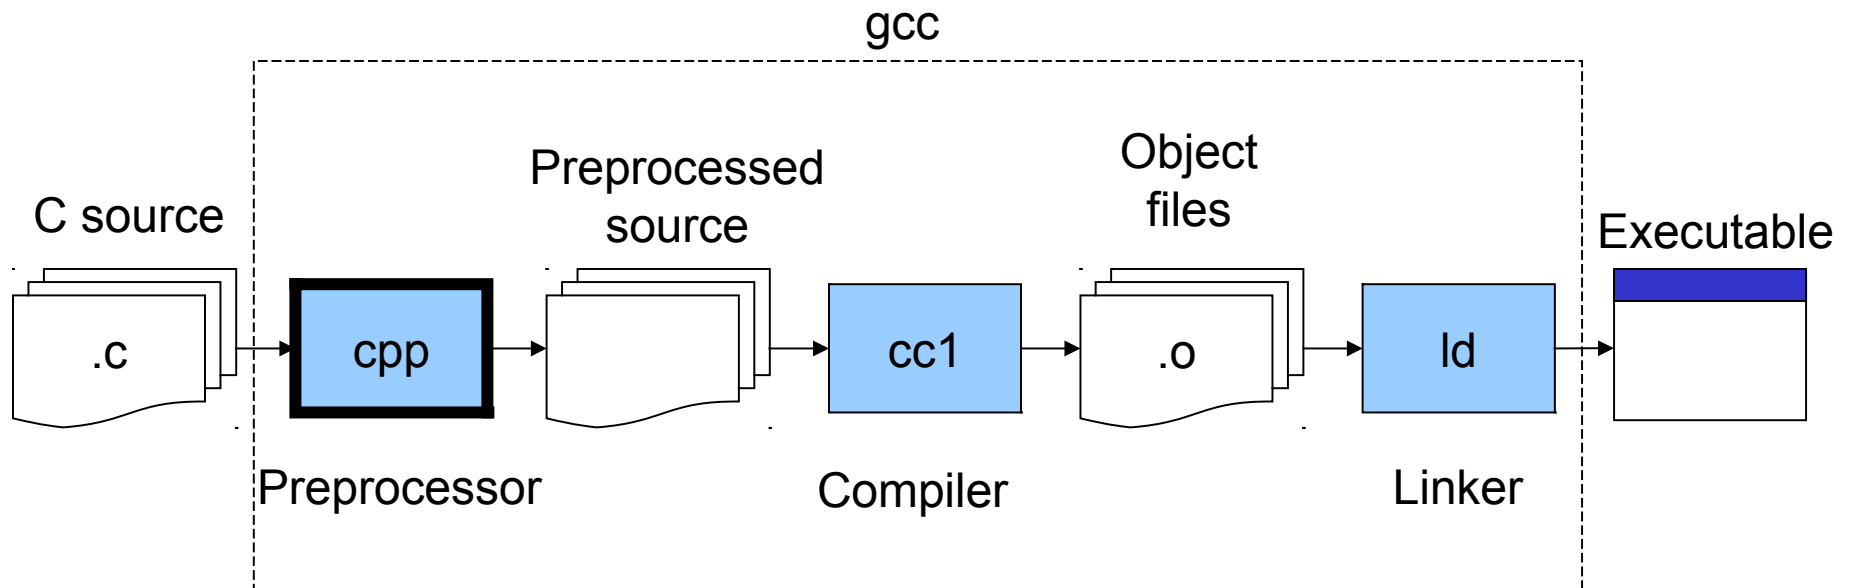
```
>> gcc ./main.c ./foo.c
>> ./a.out
global=10
```

- How did global in foo.c find global in main.c?
- How did foo() in main.c find foo() in foo.c?
- Where is the code for printf() and how does the call to printf() find it?
- What does a.out look like in the file system? How does it look like while executing in memory?

Compiler



Preprocessor



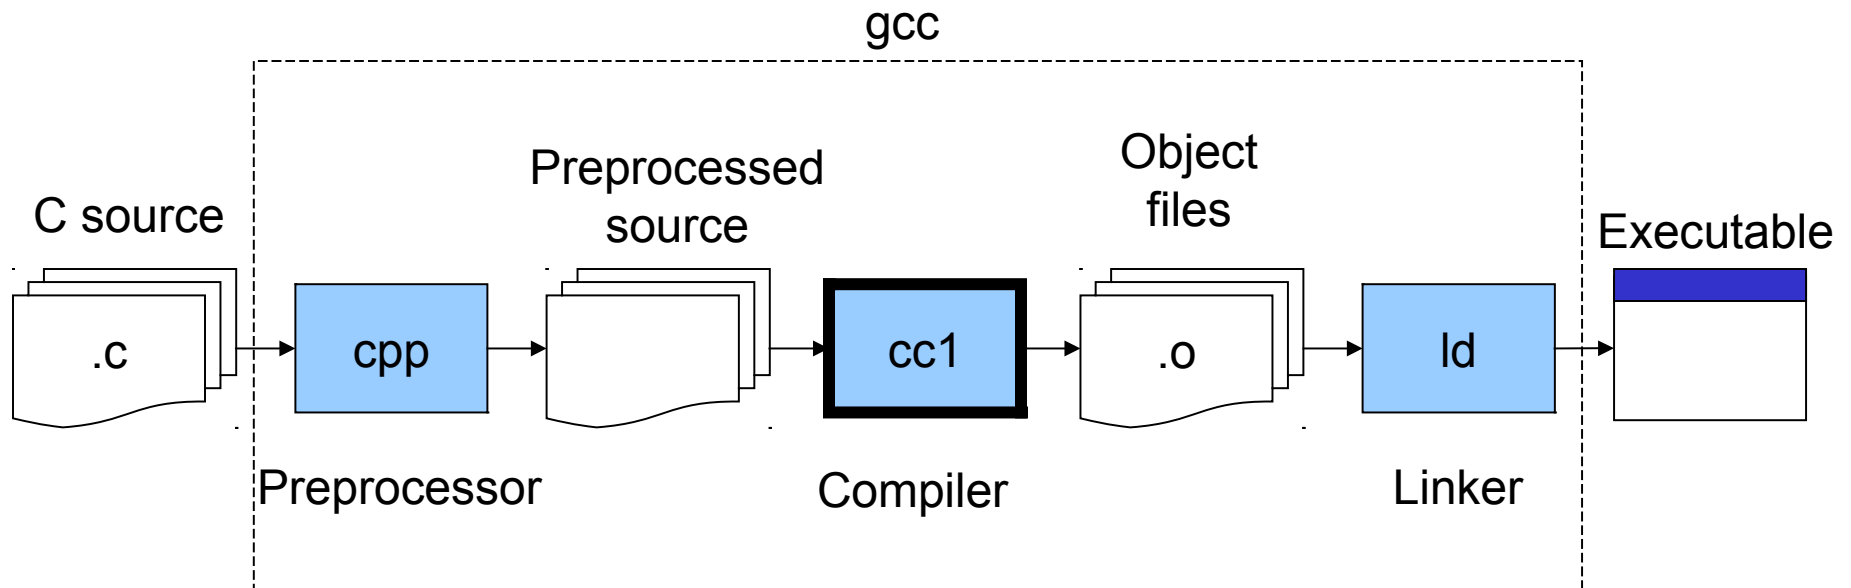
Preprocessor

- Input: A C source file with directives
- Output: A C source file after directives have been processed and comments stripped off.
- `#include` directive
 - `#include <...>` or `#include "..."`
 - Copies the contents of the file to source file
 - Contain function/variable declarations and type definitions
 - `<...>` searches for file under standard include path (usually `/usr/include`)
 - `"..."` searches for file under local directory or paths given as `-I` arguments (e.g. `gcc -I ~/local/include main.c`)

Preprocessor

- `#define` directive
 - Defines a macro (a symbol that is expanded to some other text)
 - `#define PI 3.14`
 - All instances of `PI` gets replaced by `3.14` in source
 - `#define AVG(a, b) ((a) + (b)) / 2)`
 - `AVG(2, 4)` in source gets replaced by `((2 + 4) / 2)`
 - No type checking
 - CPP is just a text translator; no concept of types
 - Will even work on `AVG("Hello", "World")`

Compiler

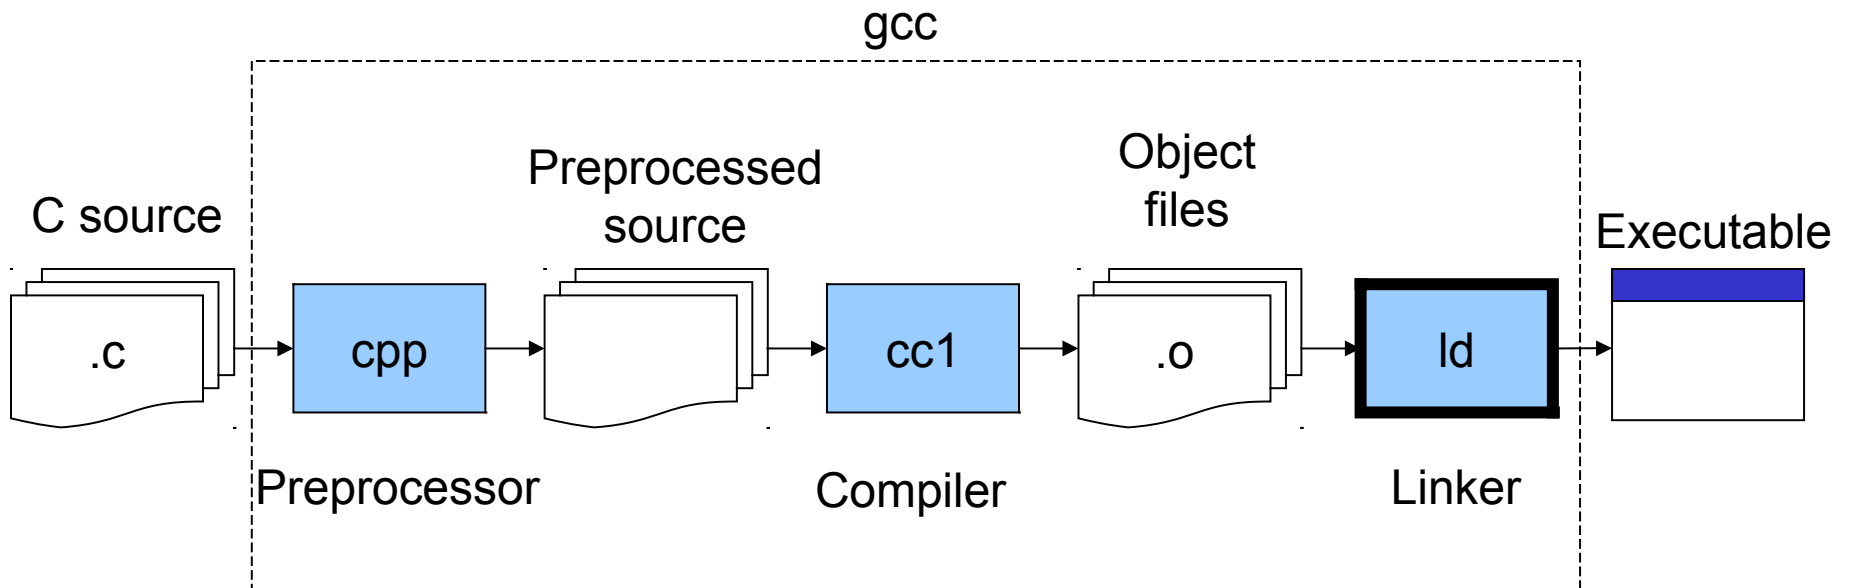


Compiler

- Input: A preprocessed C source file
- Output: An object file in machine language
- In the process it performs:
 - Syntax and type checking
 - Compiler optimizations to reduce execution time and memory footprint
 - E.g. `gcc -O2 main.c` (applies “level 2” optimizations)
 - Code generation for given machine
- One object file for one C source file
- The `-c` option produces an object file instead of an executable
- Regular `gcc main.c` command translates to the following:

```
gcc -c main.c (produces main.o)
gcc main.o (produces a.out)
rm main.o
```


Linker



Why Linkers?

- Linkers allow the combination of multiple files into executables
 - object files and libraries (libraries are archives of object files)
- Modularity
 - Large program can be written as a collection of smaller files, rather than one monolithic mass
 - Can build libraries of common functions
 - e.g., Math library, C standard library
- Efficiency
 - Time:
 - Change one source file, compile, and then re-link
 - No need to recompile other source files
 - Space:
 - Libraries of common functions can be put in a single file and linked to multiple programs

What does a Linker do?

- Step 1: Symbol resolution

- Programs define and reference symbols (variables and functions)

```
void foo() {...} /* define symbol foo */  
foo(); /* reference symbol foo */  
int *p = &x; /* define p, reference x */
```

- Symbol definitions are stored (by compilers) in a symbol table
 - Symbol table is an array of structs
 - Each entry includes name, type, size, and location of symbol
- Linker associates each symbol reference with exactly one symbol definition
 - Local symbols or static global symbols resolved within single object file
 - Only symbols with external linkage need to be in symbol table

What does a Linker do?

- Step 2: Relocation
 - Merges separate code and data sections into single sections
 - Relocates symbols from their relative locations in the .o files to their final absolute memory locations in the executable
 - Updates all references to these symbols to reflect their new positions using relocation tables
- Think of linkers as binary rewriters
 - Uses symbol tables and relocation tables to rewrite addresses while mapping object files to memory space
 - Treats the rest of the object file as a black box

Executables

- Includes everything needed to run on a system
 - Code
 - Data
 - Dependencies
 - Directions for laying out program in memory
- Self contained files that adhere to a standard format
 - Generating executables is the job of the linker

Older Executable Formats

- a.out (Assembler OUTput)
 - Oldest UNIX format
 - No longer commonly used
- COFF (Common Object File Format)
 - Older UNIX Format
 - No longer commonly used

Modern Executable Formats

- PE (Portable Executable)
 - Based on COFF
 - Used in 32- and 64-bit Windows
- ELF (Executable and Linkable Format)
 - Linux/UNIX
- Mach-O file
 - Mac

Header

- Every **a.out** formatted binary file begins with an exec structure:

```
struct exec {  
    unsigned long    a_midmag; //magic number  
    unsigned long    a_text;  // size of text segment  
    unsigned long    a_data;  // size of initialized data  
    unsigned long    a_bss;   // size of uninitialized data  
    unsigned long    a_syms;  // size of symbol table  
    unsigned long    a_entry; // entry point of program  
    unsigned long    a_trsize; // size of text relocation  
    unsigned long    a_drsize; // size of data relocation  
};
```


Contents of an Executable

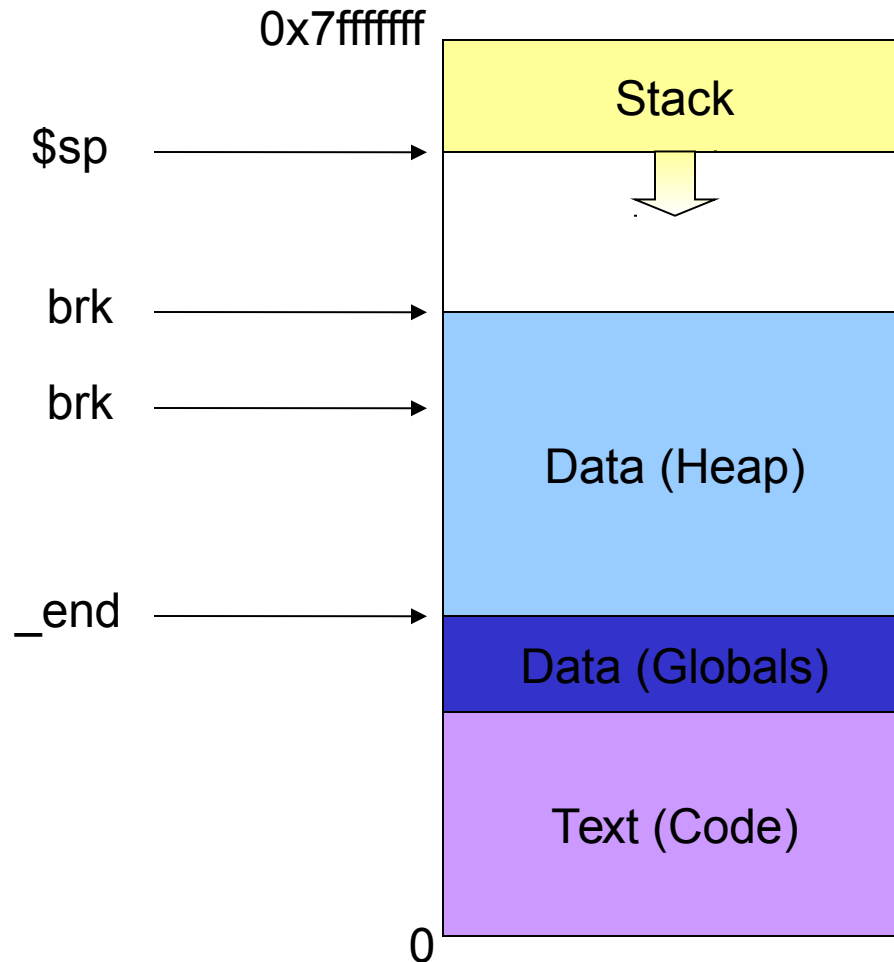
- exec header
- text segment
- data segment
- text relocations
- data relocations
- symbol table
- string table

Header

- Every **a.out** formatted binary file begins with an exec structure:

```
struct exec {  
    unsigned long    a_midmag; //magic number  
    unsigned long    a_text;  // size of text segment  
    unsigned long    a_data;  // size of initialized data  
    unsigned long    a_bss;   // size of uninitialized data  
    unsigned long    a_syms;  // size of symbol table  
    unsigned long    a_entry; // entry point of program  
    unsigned long    a_trsize; // size of text relocation  
    unsigned long    a_drsize; // size of data relocation  
};
```

Process's Address Space



Symbol Tables

```
#include <stdio.h>
int global = 0;
void foo();
int main()
{
    foo();
    printf("global=%d\n", global);
    return 0;
}
```

<main.c>

```
extern int global;
void foo() { global += 10; }
```

<foo.c>

```
>> gcc -c ./main.c ./foo.c
```

```
>> nm ./main.o
```

```
U foo
```

```
00000000 B global
```

```
00000000 T main
```

```
U printf
```

```
>> nm ./foo.o
```

```
00000000 T foo
```

```
U global
```

- **nm** is a utility that shows the symbol table of object files and executables
- offset + symbol type + symbol name
- Symbol types: undefined (**U**), text (**T**), BSS (**B**).
- Symbols *foo* and *printf* are undefined (**U**) in main.o
- Symbol *global* is undefined (**U**) in foo.o

Symbol Tables

```
#include <stdio.h>
int global = 0;
void foo();
int main()
{
    foo();
    printf("global=%d\n", global);
    return 0;
}
```

<main.c>

```
extern int global;
void foo() { global += 10; }
```

<foo.c>

```
>> gcc -c ./main.c ./foo.c
>> nm ./a.out
[sic]
080483f0 T foo
08049684 B global
               U printf@@GLIBC_2.0
```

- Symbols *foo*, *global* defined and point to appropriate places in memory space
- Symbol *printf* still undefined
 - Will be defined at load time

Relocation Tables

```
#include <stdio.h>
int global = 0;
void foo();
int main()
{
    foo();
    printf("global=%d\n", global);
    return 0;
}
```

<main.c>

```
extern int global;
void foo() { global += 10; }
```

<foo.c>

```
>> gcc -c ./main.c ./foo.c
```

```
>> objdump -r ./main.o
```

[sic]

OFFSET	TYPE	VALUE
0000000a	R_386_PC32	foo
00000010	R_386_32	global
00000015	R_386_32	.rodata
00000021	R_386_PC32	printf

- Offsets store location of reference to symbol
- For example, in main.o call to foo() will be rewritten by the linker at the relocation phase.

Relocation Tables

```
#include <stdio.h>
int global = 0;
void foo();
int main()
{
    foo();
    printf("global=%d\n", global);
    return 0;
}
```

<main.c>

```
extern int global;
void foo() { global += 10; }
```

<foo.c>

```
>> gcc -c ./main.c ./foo.c
```

```
>> objdump -R ./a.out
```

[sic]

OFFSET	TYPE	VALUE
--------	------	-------

[sic]

08049674	R_386_JUMP_SLOT	printf
----------	-----------------	--------

- Relocation table of a.out contains printf since it has not yet been resolved (no printf definition in the source files).

Libraries

- Not all code in a program is what you wrote (e.g. printf in C standard library).
- Code written by others can be included in your own program.
- How?


Linking

- Static Linking
 - Copy code into executable at compile time
 - Done by linker
- Dynamic Linking
 - Copy code into Address Space at load time or later
 - Done by link loader

Static Linking

```
#include <stdio.h>
#include <math.h>

int main() {
    printf("The sqrt of 9 is %f\n", sqrt(9));
    return 0;
}
```



Static Linking

```
#include <stdio.h>
int global = 0;
void foo();
int main()
{
    foo();
    printf("global=%d\n", global);
    return 0;
}
```

<main.c>


```
extern int global;
void foo() { global += 10; }
```

<foo.c>

```
>> gcc -m32 -static ./main.c ./foo.c -lc
>> ls -l ./a.out
-rwxr-xr-x 1 wahn UNKNOWN1 639385 Sep 30 15:20 ./a.out
>> objdump -R ./a.out
objdump: ./a.out: not a dynamic object
objdump: ./a.out: Invalid operation
```

- -static tells gcc linker to do static linking
- -lc links in library libc.a (found in /usr/lib)
- Static linking copies over all libraries into binary (just like with object files)
 - That's why a.out is so large!

Dynamic Linking



Dynamic Linking

```
#include <stdio.h>
int global = 0;
void foo();
int main()
{
    foo();
    printf("global=%d\n", global);
    return 0;
}
```

<main.c>

```
extern int global;
void foo() { global += 10; }
```

<foo.c>

```
>> gcc -m32 ./main.c ./foo.c -lc
>> ls -l ./a.out
-rwxr-xr-x 1 wahn UNKNOWN1 4769 Sep 30 15:26 ./a.out
>> ldd ./a.out
        linux-gate.so.1 => (0x00110000)
        libc.so.6 => /lib/libc.so.6 (0x00bbd000)
        /lib/ld-linux.so.2 (0x00b9b000)
```


- By default gcc does dynamic linking
- Now ./a.out is much smaller!
- -lc marks dependency on libc.so
- ldd prints out shared library dependencies
 - Program will not start if not satisfied

Dynamic Linking Pros/Cons

- Pros
 - More efficient use of storage
 - More efficient use of memory
 - One copy of library can be mapped to multiple processes
 - Much easier to fix bugs – don't have to statically relink every application
- Cons
 - Address resolution done at load time – slow
 - Versioning of multiple dynamic libraries (DLL Hell)
 - Security holes due to replacement of libraries

Dynamic Loading

and Dynamic Link Libraries



Function Pointers

- How do we call a function when we can't be sure what address it's loaded at?
- Need a level of indirection.
- Use a function pointer.

Dynamic Loading

```
#include <dlfcn.h>
int main() {
    void *handle;
    int (*pf)(const char *format, ...);
    /* open the C standard library */
    handle = dlopen("libc.so.6", RTLD_LAZY);
    /* find the address of the printf function */
    pf = dlsym(handle, "printf");
    (*pf)("Hello world!\n");
    dlclose(handle);
    return 0;
}
```

```
>> gcc -ldl ./dlsym.c
>> ./a.out
Hello world!
```

Problem 1

- You have a commodity computer. There is a text file of size 4TB, written on the 12TB HDD. The file contains short integers only. Every line of the file contains exactly one number. The numbers are different.
- Propose a solution for sorting the file with minimal number of operations.

Problem 2

- Implement the command **wc**. When run, the program should get 2 arguments. The first one is an option and the second one is a file name: *wc <option> <file_name>*.
- Option '**c**' instructs the program to count the number of characters. Option '**w**' instructs it to count the number of words and option '**l**' – the number of lines.

Problem 3

- Write a function in C, which gets a string as input parameter and checks if it is a palindrome. The function should return 1, if the string is a palindrome and -1 otherwise.
- A **palindrome** is a word, phrase, number, or other sequence of characters which reads the same backward or forward.