


System Calls & Signals

CS449 Spring 2016

Operating system

- OS – a layer of software interposed between the application program and the hardware

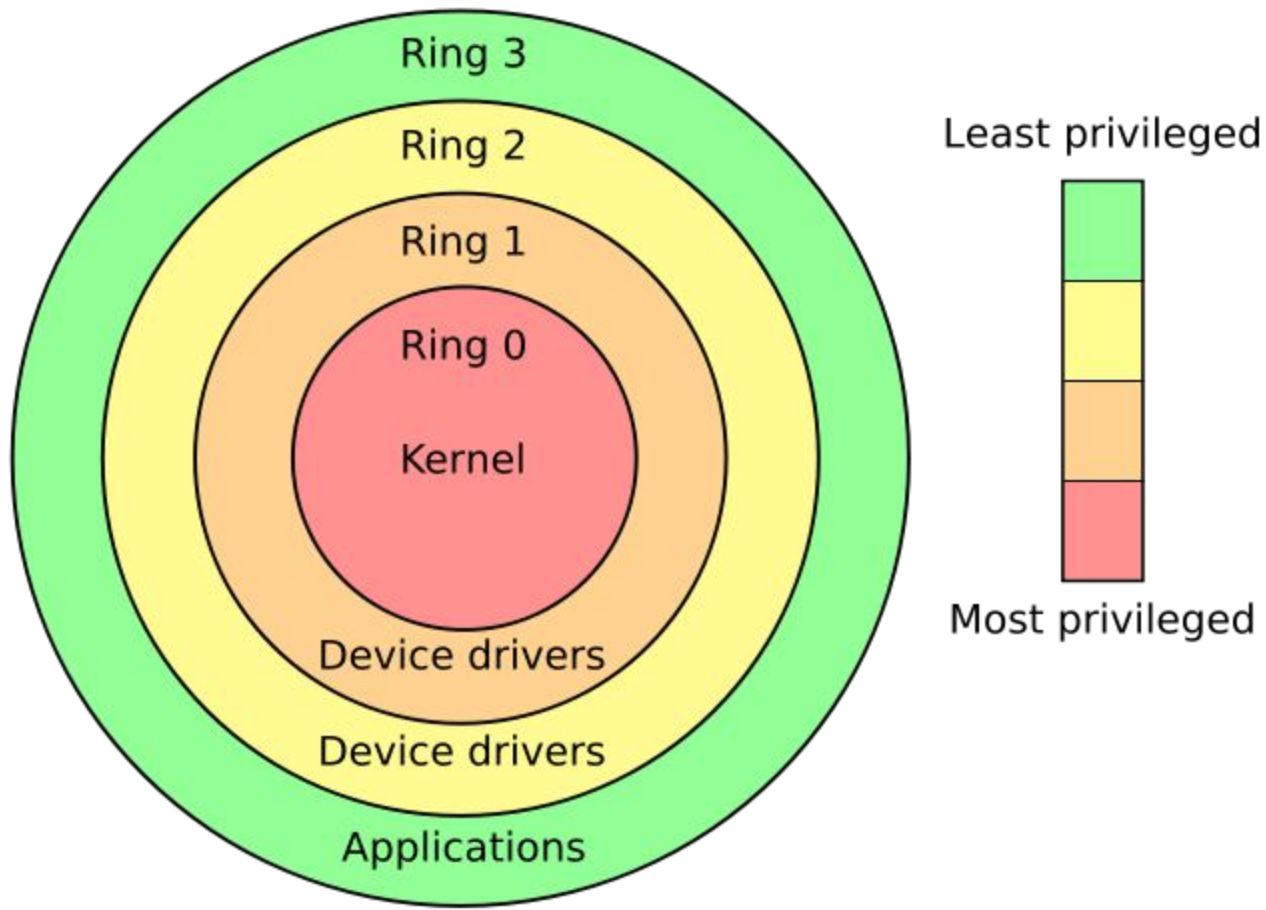


- Two primary goal
 - Manage hardware resources
 - Abstract details about hardware
 - I/O devices abstracted by files
 - Main memory abstracted by virtual memory
 - Processors abstracted by processes

What is an OS

- Pedantically it is just a **Kernel**
 - A program that implements the core OS functions (process scheduling, memory/HW management)
- More commonly it refers to the basic environment
 - Libraries, user interface, foundational programs
- A Kernel is a program, but it is not an application
 - Doesn't actually execute on its own (except at boot time)
 - **Only reacts to events and requests**
 - System calls, exceptions, and interrupts

x86 Privilege Levels



OS Interactions

- **System Calls**

- Allow an application to request the OS to do something
 - Read a file, print to console, etc...

- **Signals**

- OS notifies an application of something
 - SEGFAULT, Killed, etc...

- **Exceptions**

- Application does something it wasn't supposed to
 - Access invalid memory address

- **Interrupts**

- Processor receives a hardware signal on an interrupt pin
 - Mouse move or keyboard type

Signals

- Notifications sent to a program by OS
 - Indicate special events
- Allows for asynchronous notification rather than polling
- Polling – to explicitly ask if something occurred, usually repeatedly

kill -l

(returns a complete list of signals for your system)

SIGHUP	SIGINT	SIGQUIT	SIGILL	SIGTRAP
SIGABRT	SIGBUS	SIGFPE	SIGKILL	SIGUSR1
SIGSEGV	SIGUSR2	SIGPIPE	SIGALRM	SIGTERM
SIGCHLD	SIGCONT	SIGSTOP	SIGTSTP	SIGTTIN
SIGTTOU	SIGURG	SIGXCPU	SIGXFSZ	SIGVTALRM
SIGPROF	SIGWINCH	SIGIO	SIGPWR	SIGSYS
SIGRTMIN	SIGRTMIN+1	SIGRTMIN+2	SIGRTMIN+3	SIGRTMIN+4
SIGRTMIN+5	SIGRTMIN+6	SIGRTMIN+7	SIGRTMIN+8	SIGRTMIN+9
SIGRTMIN+10	SIGRTMIN+11	SIGRTMIN+12	SIGRTMIN+13	SIGRTMIN+14
SIGRTMIN+15	SIGRTMAX-14	SIGRTMAX-13	SIGRTMAX-12	SIGRTMAX-11
SIGRTMAX-10	SIGRTMAX-9	SIGRTMAX-8	SIGRTMAX-7	SIGRTMAX-6
SIGRTMAX-5	SIGRTMAX-4	SIGRTMAX-3	SIGRTMAX-2	SIGRTMAX-1
SIGRTMAX				

Common Error Signals

- **SIGILL** – Illegal Instruction
- **SIGBUS** – Bus Error, usually caused by bad data alignment or a bad address
- **SIGFPE** – Floating Point Exception
- **SIGSEGV** – Segmentation violation, i.e., a bad address

Termination Signals

- **SIGINT** – Interrupt, or what happens when you hit CTRL + C
- **SIGTERM** – Ask nicely for a program to end (can be caught)
- **SIGKILL** – Ask meanly for a program to end (cannot be caught)
- **SIGABRT, SIGQUIT** – End a program with a core dump

kill

- From the shell in UNIX you can send signals to a program.
- Use ps to get a process ID

```
(1) thot $ ps -af
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
jrmst106	27500	27470	0	07:13	???	00:00:00	crashed_program
jrmst106	27507	27474	0	07:13	pts/5	00:00:00	ps -af

- kill it!

```
kill 27500 – Sends SIGTERM
```

```
kill -9 27500 – Sends SIGKILL
```

kill

- `kill()` is the system call that can send a process a signal (any signal, not just `SIGKILL`)

```
#include <unistd.h>
#include <sys/types.h>
#include <signal.h>

int main() {
    pid_t my_pid = getpid();
    kill(my_pid, SIGSTOP);
    return 0;
}
```

Catching Signals

- Most signals can be caught
 - Like exceptions in Java
- Do some cleanup, then exit
- Generally bad to try to continue, the machine might be in a corrupt state
- Some signals can be caught safely though

SIGALRM

```
#include <unistd.h>
#include <signal.h>

int timer = 10;

void catch_alarm(int sig_num) {
    printf("%d\n", timer--);
    alarm(1);
}

int main() {
    signal(SIGALRM, catch_alarm);

    alarm(1);
    while(timer > 0) ;
    alarm(0);
    return 0;
}
```

```
>> ./a.out
```

```
10
```

```
9
```

```
8
```

```
7
```

```
6
```

```
5
```

```
4
```

```
3
```

```
2
```

```
1
```

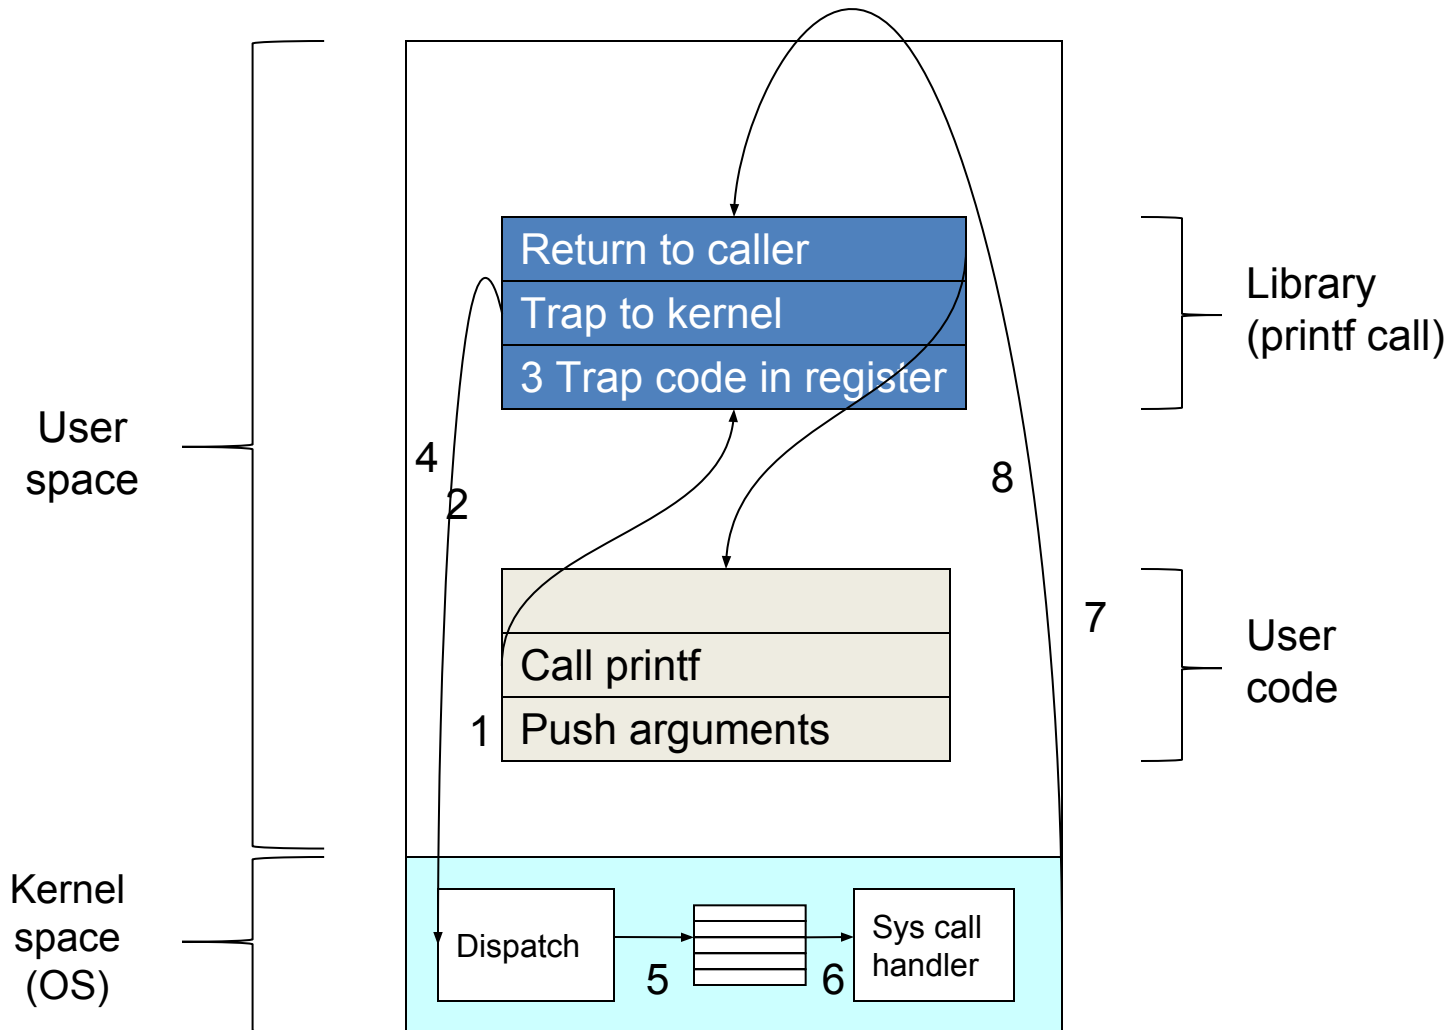
Exceptions

- You've seen these in Java
 - Exceptions are generated by unexpected errors
- Hardware exceptions exist for the same purpose
 - Hardware has to tell OS that a program tried to do something illegal
 - Executed an illegal instruction, violated memory protection
 - OS then has to respond to the violation
 - Handle the exception let program continue to run (e.g. page fault)
 - Page Fault: exception that happens when accessing page in disk
 - Notify the program that something happened (via signal)
 - Kill the program

Exception Handling

- Each exception has an entry in a dispatch table
 - An array of function pointers to specific handlers
 - Interrupt Vector Table
 - Appropriate privilege given on jump to handler
 - Registers must be saved to/restored from kernel stack
- Hardware interrupts handled by same dispatch table
- System calls often implemented using exceptions
 - x86 versions uses a trap instruction: int 0x80
 - Trap: an intentional exception to invoke a specific handler
 - Jumps to entry 0x80 in dispatch table

System Call



strace ./hello

```
fstat(1, {st_mode=S_IFCHR|0600, st_rdev=makedev(136, 7), ...}) = 0
```

```
mmap(NULL, 4096, PROT_READ|PROT_WRITE,  
MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) =  
0x2a95557000
```

```
write(1, "Hello world!\n", 13Hello world!) = 13
```

```
exit_group(0)
```

Linux Syscalls

- 325 syscall slots reserved in system call table (2.6.23.1 kernel) -- Not all are used

exit	Causes a process to terminate
fork	Creates a new process, identical to the current one
read	Reads data from a file or device
write	Writes data to a file or device
open	Opens a file
close	Closes a file
creat	Creates a file

Using Syscalls

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int main()
{
    int fd;
    char buffer[100];
    strcpy(buffer, "Hello, World!\n");

    fd = open("hello.txt", O_WRONLY | O_CREAT);
    write(fd, buffer, strlen(buffer));
    close(fd);
    exit(0);

    return 0;
}
```

OR-ing Flags

- Define constants as powers of 2
- Bitwise OR to combine
- Bitwise AND to test

```
#define O_RDONLY      0
#define O_WRONLY      1
#define O_RDWR        2
#define O_CREAT       16
```

File Descriptors

- Integer identifying a unique open file
 - Similar to FILE *
 - FILE struct has a file descriptor member
- File descriptors index into a file descriptor table
- OS maintains additional information about the file to do things such as clean up on process termination
- Three standard file descriptors opened automatically:
 - 0 – stdin
 - 1 – stdout
 - 2 – stderr

Dumping the File Descriptor Table

- `ls -l /proc/self/fd` (listing default file descriptors)

```
(84) thot $ ls -l /proc/self/fd
```

```
total 0
```

```
lrwx----- 1 wahn UNKNOWN1 64 Sep 11 13:37 0 -> /dev/pts/0
```

```
lrwx----- 1 wahn UNKNOWN1 64 Sep 11 13:37 1 -> /dev/pts/0
```

```
lrwx----- 1 wahn UNKNOWN1 64 Sep 11 13:37 2 -> /dev/pts/0
```

```
lr-x----- 1 wahn UNKNOWN1 64 Sep 11 13:37 3 -> /proc/16970/fd
```

- `ls -l /proc/self/fd > /tmp/ls.out` (listing with redirect)

```
(85) thot $ ls -l /proc/self/fd > /tmp/ls.out && cat /tmp/ls.out
```

```
total 0
```

```
lrwx----- 1 wahn UNKNOWN1 64 Sep 11 13:41 0 -> /dev/pts/0
```

```
l-wx----- 1 wahn UNKNOWN1 64 Sep 11 13:41 1 -> /tmp/ls.out
```

```
lrwx----- 1 wahn UNKNOWN1 64 Sep 11 13:41 2 -> /dev/pts/0
```

```
lr-x----- 1 wahn UNKNOWN1 64 Sep 11 13:41 3 -> /proc/17025/fd
```