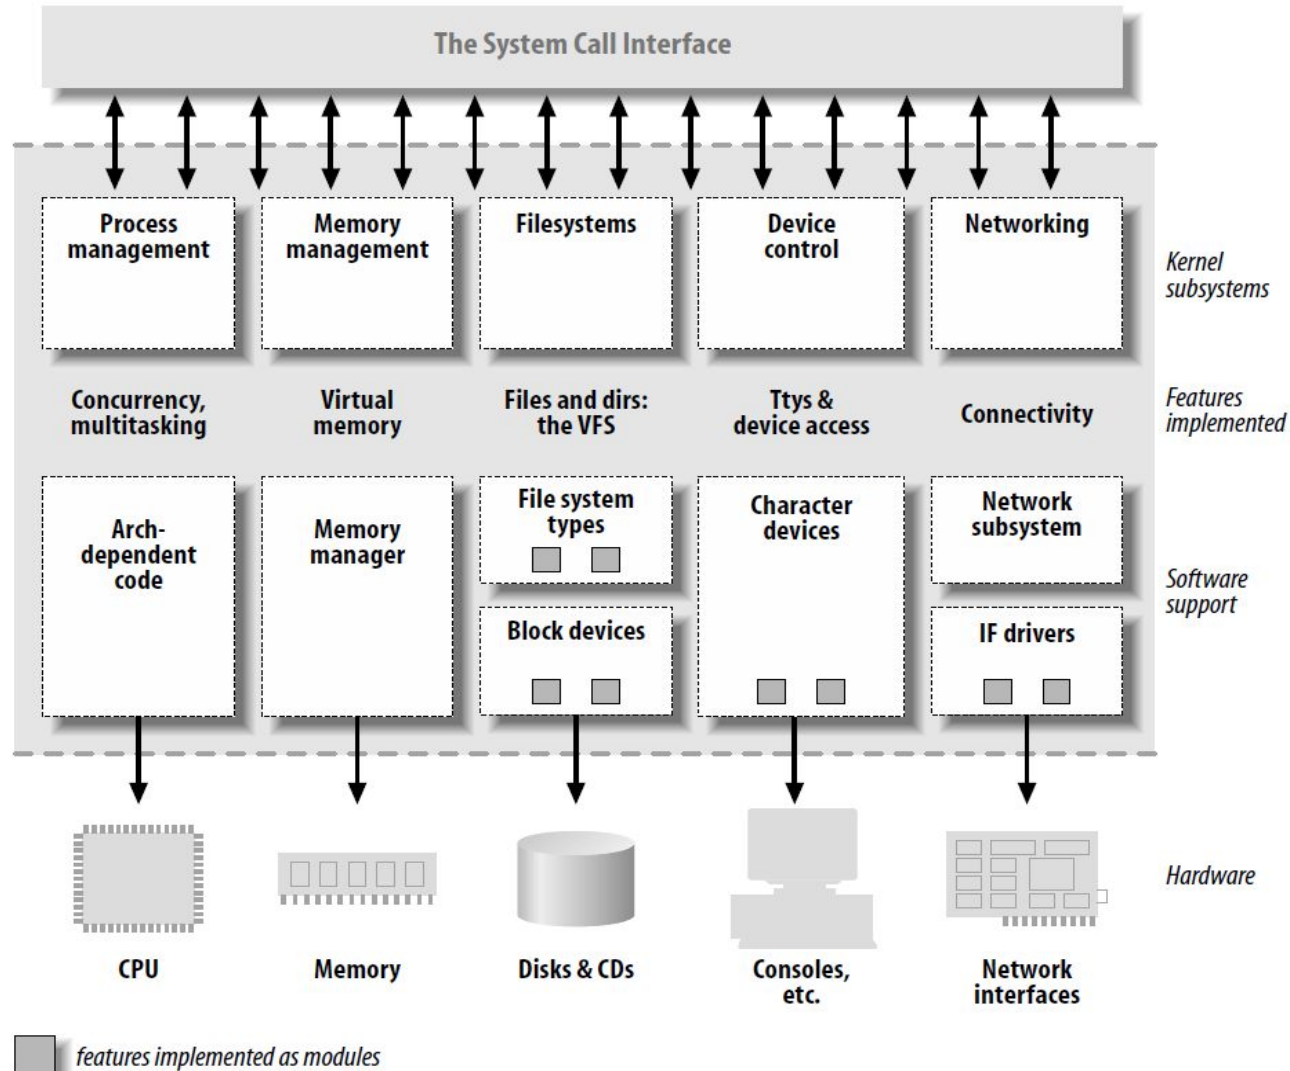


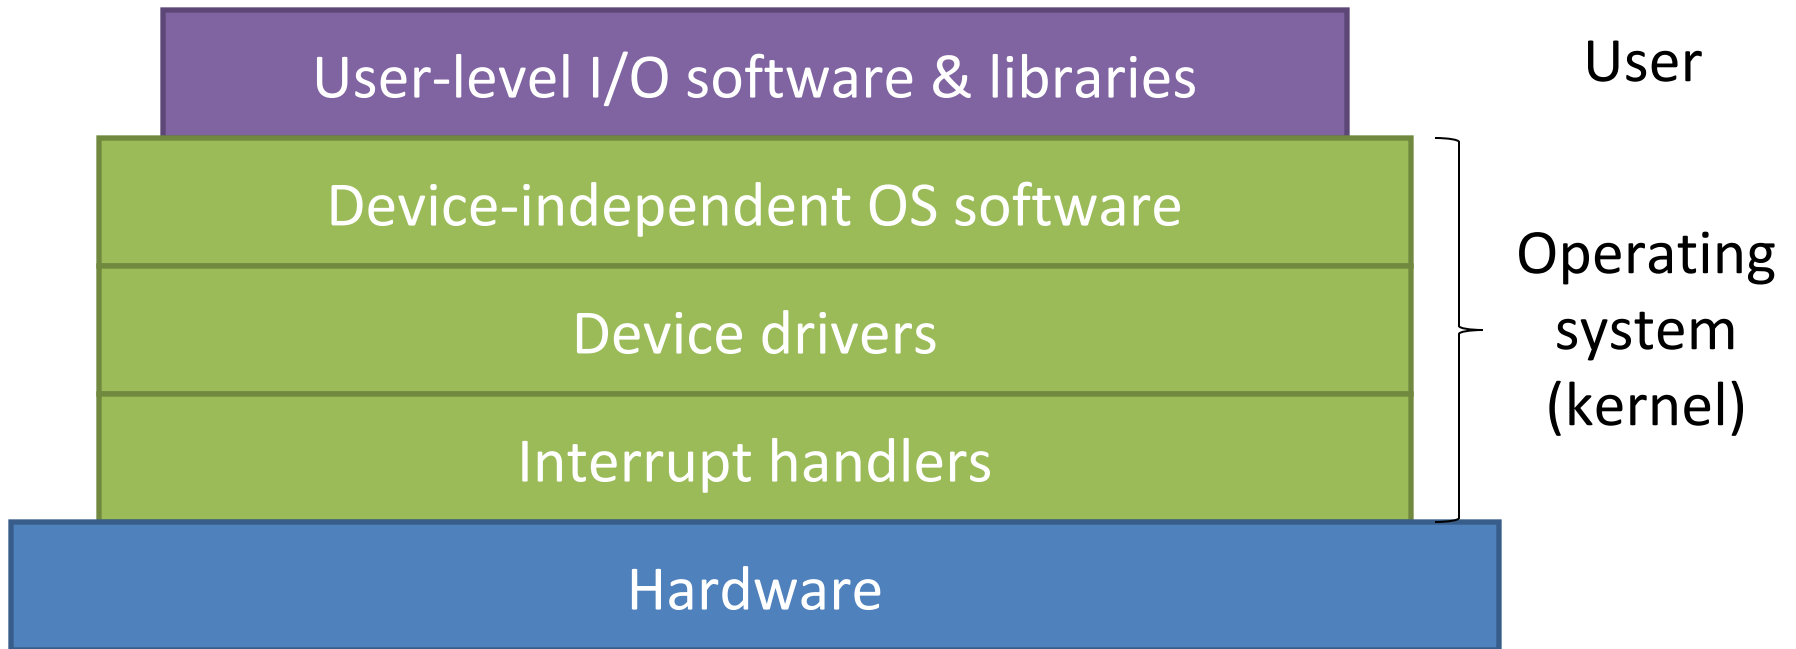
# Device Drivers

CS449 Spring 2016

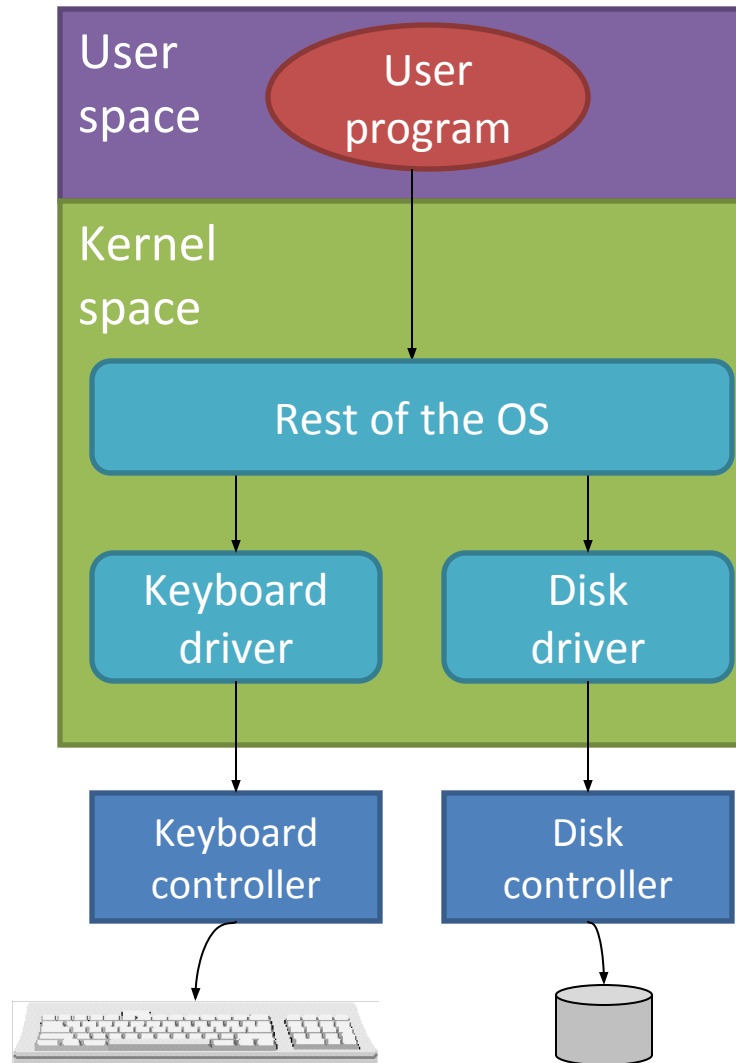
# Abstraction via the OS



# Software Layers



# Device Drivers



# Types of Devices

- **Block Devices**

- A device that stores data in fixed-sized blocks, each uniquely addressed, and can be randomly accessed
- E.g., Disks, Flash Drives

- **Character Devices**

- Device that delivers or accepts a stream of characters
- E.g., Keyboard, mouse, terminal

# Mechanism vs. Policy

- Mechanism – What capabilities are provided
  - E.g. Hard disk driver: exposes the disk as a continuous array of data blocks
- Policy – How to use those capabilities
  - E.g. Data blocks can be organized using a file system, or they can be used as a raw block device
  - E.g. Data blocks may only be accessible by certain users
- Drivers should be flexible by only providing mechanisms not policies

# Sysfs

- Mounted on `/sys/`
- Contains files that provide information about devices: whether it's powered on, the vendor name and model, what bus the device is plugged into, etc. It's of interest to applications that manage devices
- Exports information about devices and drivers to userspace
- Can configure aspects of device

# Devfs

- Mounted on `/dev/`
- Contains files that allow programs to access the devices themselves: write data to a serial port, read a hard disk, etc. It's of interest to applications that access devices
- Character and block devices exposed via the filesystem
- `/dev/` typically contains “files” that represent the different devices on a system
- `/dev/console` – the console
- `/dev/fd/` - a process's open file descriptors



# Device Drivers in Linux

- Can be compiled into the kernel
- Can be loaded dynamically as Modules

# Hello World Module

```
#include <linux/init.h>
#include <linux/module.h>
MODULE_LICENSE("Dual BSD/GPL");
static int hello_init(void)
{
    printk(KERN_ALERT "Hello, world\n");
    return 0;
}
static void hello_exit(void)
{
    printk(KERN_ALERT "Goodbye, cruel world\n");
}
module_init(hello_init);
module_exit(hello_exit);
```

# Why printk?

- The kernel does not have access to libraries
- Can't use printf or many other standard functions (FILE stuff, strtok, etc.)
- Modules are linked against the kernel only
- Kernel provides useful set of common functions like strcpy, strcat, etc.

# MODULE\_LICENSE

- Informs the kernel what license the module source code is under
- Affects which symbols (functions, variables, etc.) it may access in the kernel
- A GPL-licensed module can access everything
- Certain (or not specifying one) module license will “taint” the kernel

# Building & Running

% **make**

```
make[1]: Entering directory `/usr/src/linux-2.6.10'
  CC [M] /home/ldd3/src/misc-modules/hello.o
  Building modules, stage 2.
  MODPOST
  CC /home/ldd3/src/misc-modules/hello.mod.o
  LD [M] /home/ldd3/src/misc-modules/hello.ko
make[1]: Leaving directory `/usr/src/linux-2.6.10'
```

% **su**

```
root# insmod ./hello.ko
Hello, world
root# rmmod hello
Goodbye cruel world
root#
```

# Makefile

```
obj-m := hello_dev.o
```

```
KDIR := /u/SysLab/shared/linux-2.6.23.1
```

```
PWD := $(shell pwd)
```

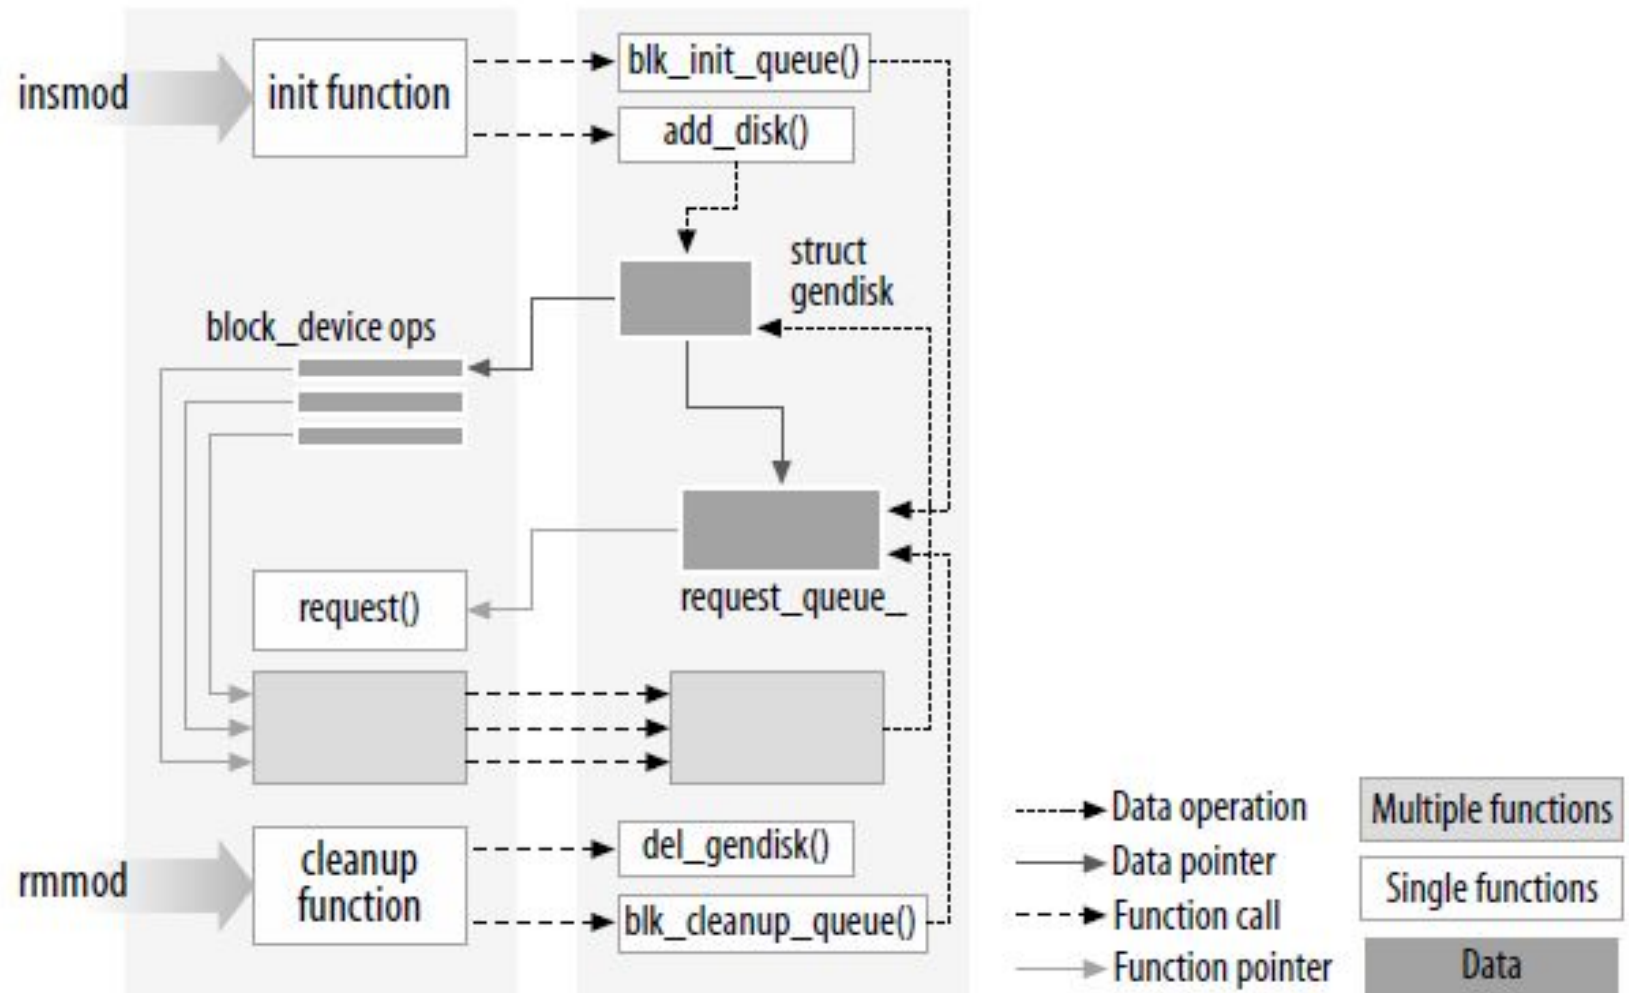
```
default:
```

```
$(MAKE) -C $(KDIR) M=$(PWD) modules
```

# Module Helper Programs

- `insmod` – loads a module
- `rmmod` – unloads a module
- `lsmod` – lists what modules are loaded
- `modprobe` – loads a module checking dependencies

# Loading & Unloading a Module





# Hello World with Read function

Download full source from:

[http://www.linuxdevcenter.com/2007/07/05/examples/hello\\_dev.c](http://www.linuxdevcenter.com/2007/07/05/examples/hello_dev.c)

```

#include <linux/fs.h>
#include <linux/init.h>
#include <linux/miscdevice.h>
#include <linux/module.h>
#include <asm/uaccess.h>

/*
 * hello_read is the function called when a process calls read() on
 * /dev/hello. It writes "Hello, world!" to the buffer passed in the
 * read() call.
 */

static ssize_t hello_read(struct file * file, char * buf, size_t count, loff_t *ppos)
{
    char *hello_str = "Hello, world!\n";
    int len = strlen(hello_str); /* Don't include the null byte. */
    /*
     * We only support reading the whole string at once.
     */
    if (count < len)
        return -EINVAL;
    /*
     * If file position is non-zero, then assume the string has
     * been read and indicate there is no more data to be read.
     */
    if (*ppos != 0)
        return 0;
    /*
     * Besides copying the string to the user provided buffer,
     * this function also checks that the user has permission to
     * write to the buffer, that it is mapped, etc.
     */
    if (copy_to_user(buf, hello_str, len))
        return -EINVAL;
    /*
     * Tell the user how much data we wrote.
     */
    *ppos = len;

    return len;
}

```

```
/*  
 * The only file operation we care about is read.  
 */
```

```
static const struct file_operations hello_fops = {  
    .owner          = THIS_MODULE,  
    .read           = hello_read,  
};
```

```
static struct miscdevice hello_dev = {  
    /*  
     * We don't care what minor number we end up with, so tell the  
     * kernel to just pick one.  
     */  
    MISC_DYNAMIC_MINOR,  
    /*  
     * Name ourselves /dev/hello.  
     */  
    "hello",  
    /*  
     * What functions to call when a program performs file  
     * operations on the device.  
     */  
    &hello_fops  
};
```

```

static int __init hello_init(void)
{
    int ret;

    /*
     * Create the "hello" device in the /sys/class/misc directory.
     * Udev will automatically create the /dev/hello device using
     * the default rules.
     */
    ret = misc_register(&hello_dev);
    if (ret)
        printk(KERN_ERR "Unable to register \"Hello, world!\" misc device\n");

    return ret;
}

module_init(hello_init);

static void __exit hello_exit(void)
{
    misc_deregister(&hello_dev);
}

module_exit(hello_exit);

MODULE_LICENSE("GPL");
MODULE_AUTHOR("Valerie Henson <val@nmt.edu>");
MODULE_DESCRIPTION("\"Hello, world!\" minimal module");
MODULE_VERSION("dev");

```

# Test the Read function

Then, we're ready to compile and load the module:

```
$ cd hello_proc
```

```
$ make
```

```
$ sudo insmod ./hello_proc.ko
```

Now, there is a file named `/proc/hello_world` that will produce "Hello, world!" when read:

```
$ cat /proc/hello_world
```

```
Hello, world!
```

# Device Operation Callbacks

```
struct file_operations {  
    struct module *owner;  
    int (*open) (struct inode *, struct file *);  
    ssize_t (*read) (struct file *, char *, size_t, loff_t *);  
    ssize_t (*write) (struct file *, const char *, size_t, loff_t *);  
    ...  
};
```

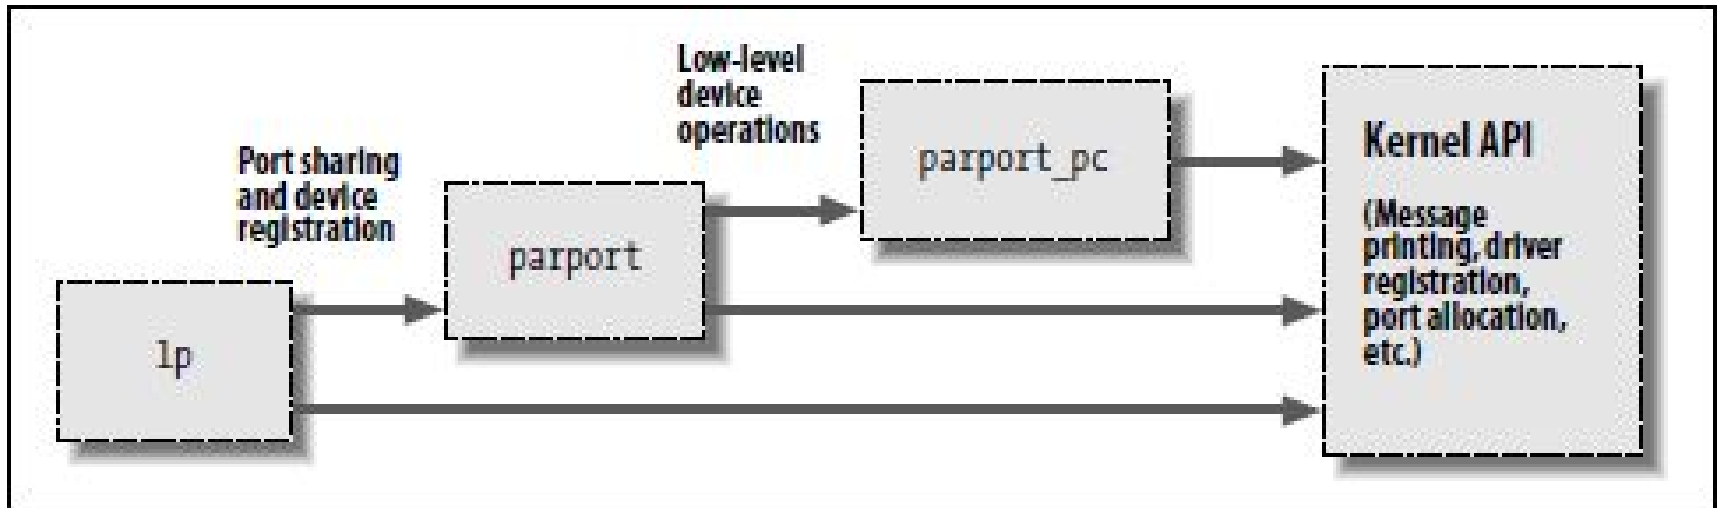
- Struct containing callbacks passed to OS when driver is loaded
- OS does callbacks on functions on corresponding system call

# Error Handling

```
int __init my_init_function(void)
{
    int err;
    /* registration takes a pointer and a name */
    err = register_this(ptr1, "driver");
    if (err) goto fail_this;
    err = register_that(ptr2, "driver");
    if (err) goto fail_that;
    err = register_those(ptr3, "driver");
    if (err) goto fail_those;
    return 0; /* success */

    fail_those: unregister_that(ptr2, "driver");
    fail_that:  unregister_this(ptr1, "driver");
    fail_this: return err; /* propagate the error */
}
```

# Driver Stacking





# Things not to do in the kernel

- Do not stack allocate big arrays
  - Kernel stack is small, maybe only a single page (4KB)
  - You should use kmalloc to allocate heap space
- Do not leave memory unfreed
  - Will stay around forever until the next reboot!
- Do not do floating point arithmetic
  - Context switch into the kernel does not save floating point registers

# Race Conditions

- The kernel will make calls into your module while your initialization function is still running
- Multiple applications will attempt to access your driver simultaneously

# User Space Drivers

- Advantages?
  - Full Standard C Library can be linked in
  - Can use a conventional debugger like GDB
  - Problems with driver will not crash entire system
  - Can be swapped to disk using virtual memory
- FUSE – Filesystem in User Space
  - Useful for implementing virtual file systems (e.g. by communicating with cloud storage)