

Threads

CS449 Spring 2016

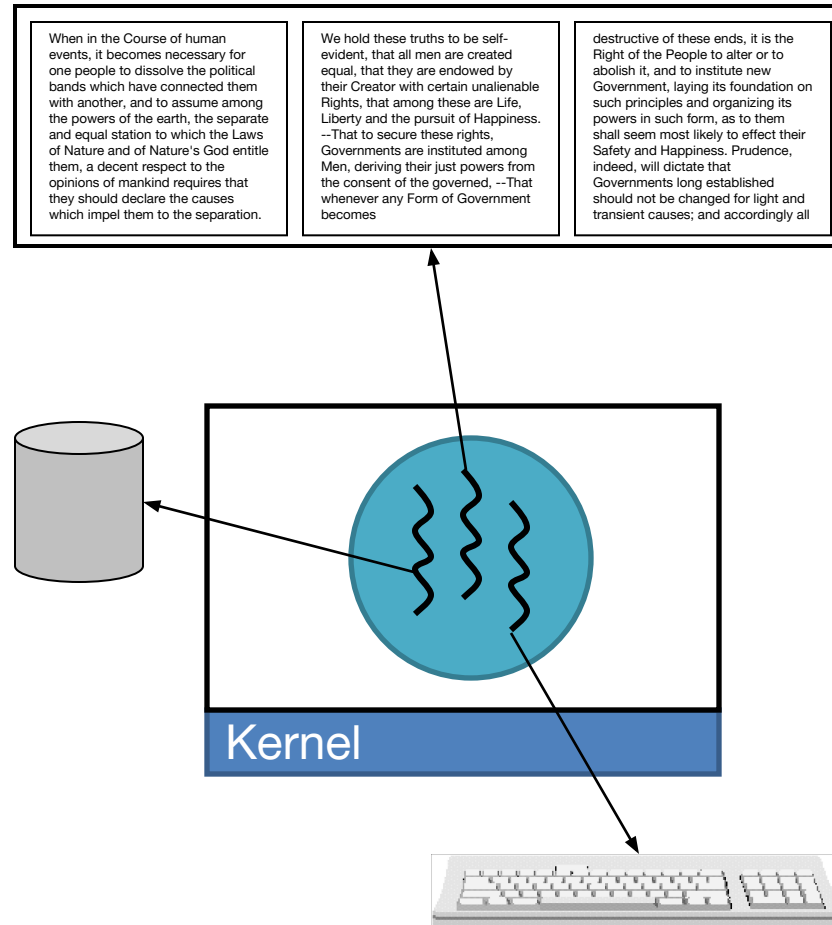
Midterm 2 Statistics

- Max: 86
- Min: 11
- Average: 55.4
- Median: 59

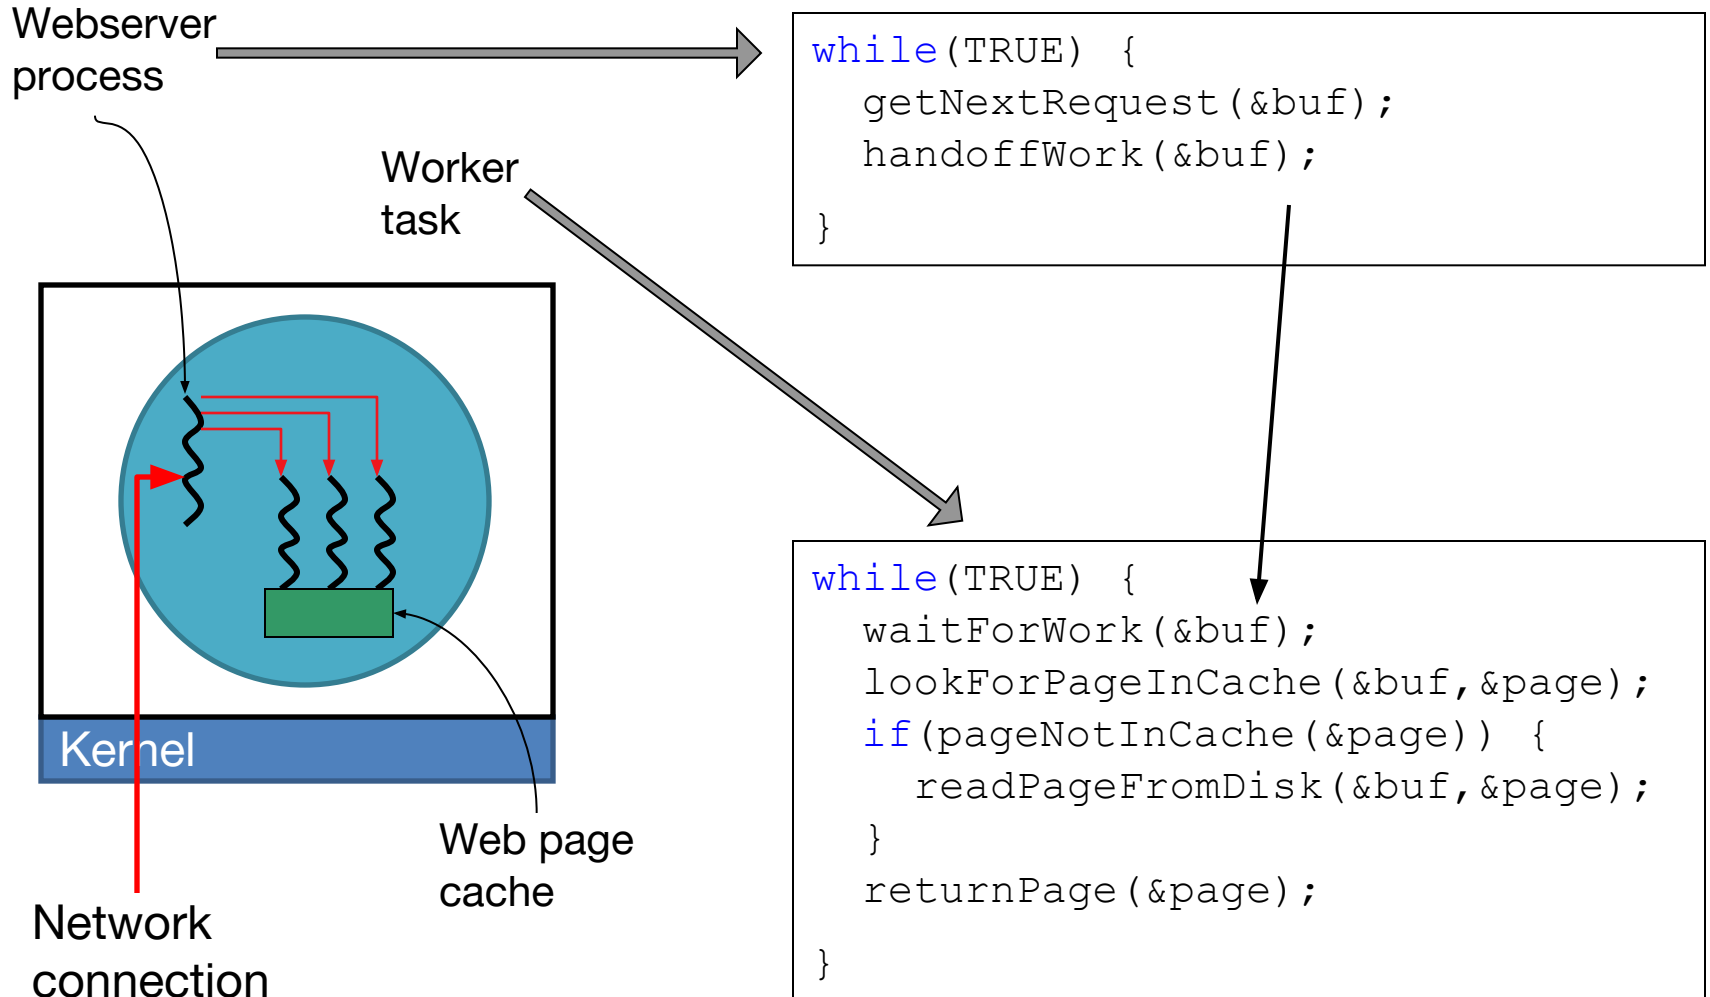
Concurrency and Parallelism

- Many programs need to perform mostly independent tasks that do not need to be serialized
 - Web server: page requests
 - Text editor: auto-save, spell checking, text entry
 - Web client: tabbed browsing
- **Parallel Programming**
 - Make multiple runtime copies of the same task
 - Underlying system runs each task on a separate CPU
 - For performance
 - Makes sense only on a multiprocessor machine
- **Concurrent Programming**
 - Divide program into multiple, generally different, tasks
 - Underlying system provides illusion of concurrent progress
 - For responsiveness (i.e. spell check does not delay text entry)
 - Can make sense even on a single processor machine
- **The need for parallelism has become more acute with multicores**

Concurrency Example (Word Processor)



Parallelism Example (Webserver)



Programs and Processes

- **Program: Executable binary (code and static data)**
- **Process: A program loaded into memory**
 - Program (executable binary with data and text section)
 - Execution state (heap, stack, and processor registers)

Program

```
int foo() {  
    return 0;  
}  
  
int main() {  
    foo();  
    return 0;  
}
```

Process

```
int foo() {  
    return 0;  
}  
  
int main() {  
    foo();  
    return 0;  
}
```

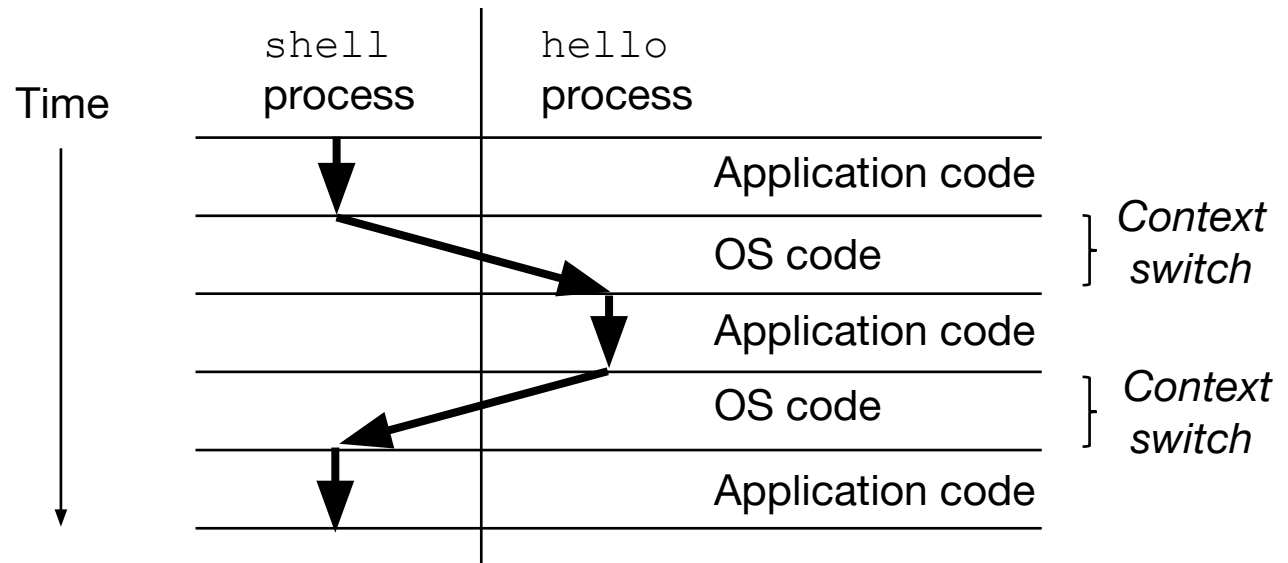
Heap

Stack

Registers


Context Switching

- OS provides the illusion of a dedicated machine per process
- Context switch
 - Saving context of one process, restoring that of another one
 - Processor alternates between executing different processes




Multiprogramming: Concurrently Running Programs

Single \$PC
(CPU's point of view)




One stream of instructions

Multiple \$PCs
(process point of view)



Isolated address spaces
that are not interrupted

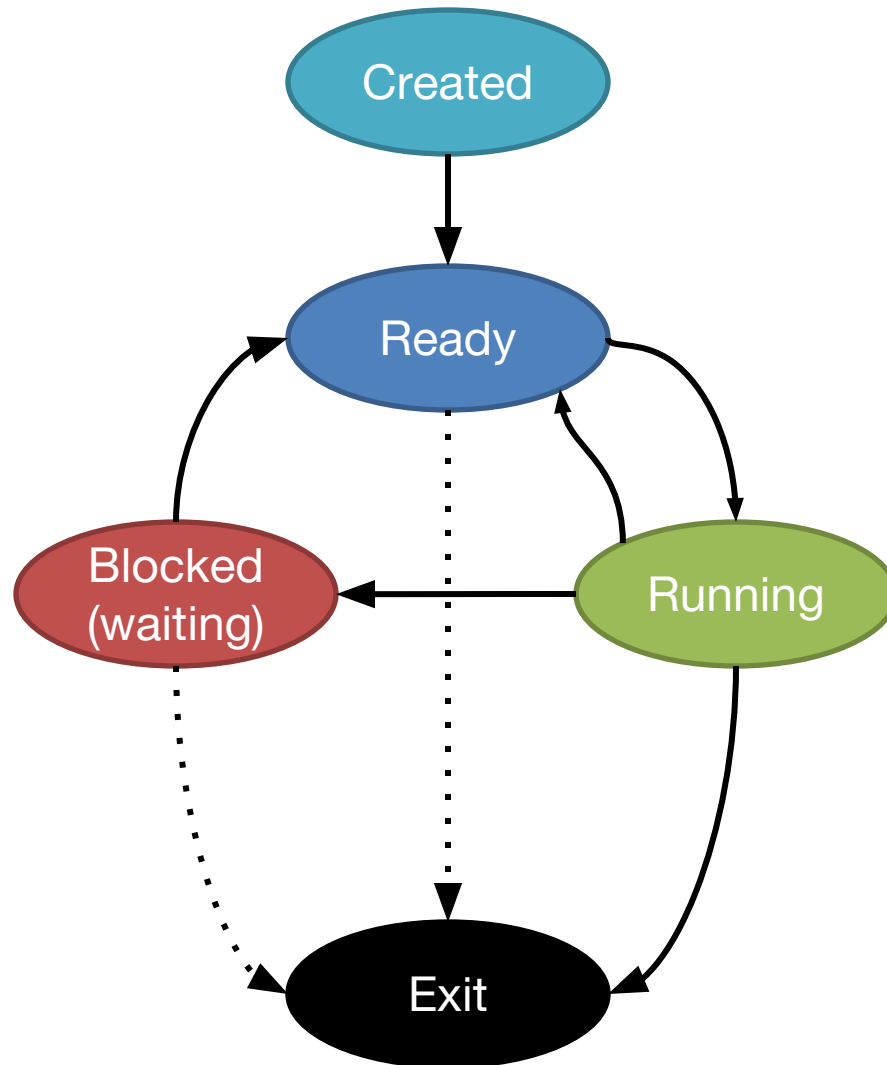


Dispatch Mechanism

- OS maintains list of all processes (PCBs)
- Each process has a mode
 - **Running**: Executing on the CPU
 - **Ready**: Waiting to execute on CPU
 - **Blocked**: Waiting for I/O or synchronization with another thread
- Dispatch Loop

```
while (1) {  
    run process for a while;  
    stop process and save its state;  
    load state of another process;  
}
```

Life Cycle of a Process



Processes Sometimes Infeasible

- Expensive in terms of memory
 - Must maintain PCB for each process in OS
 - In particular page tables
- Expensive in terms of performance
 - Forking: must duplicate PCB (including page tables)
 - Context Switching: involves system call into OS and flushing cached page table entries in MMU
 - Inter-process Communication: involves system call to invoke OS to move data
- When memory needs to be *shared* instead of *duplicated*:
 - When there is a shared modified data structure

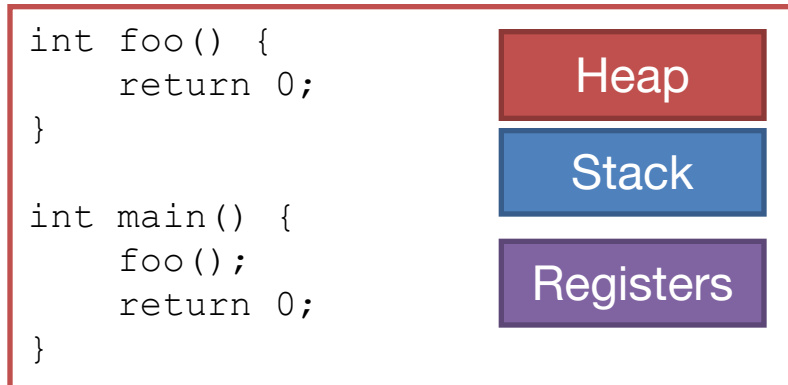
Thread

A stream of instructions and their associated state

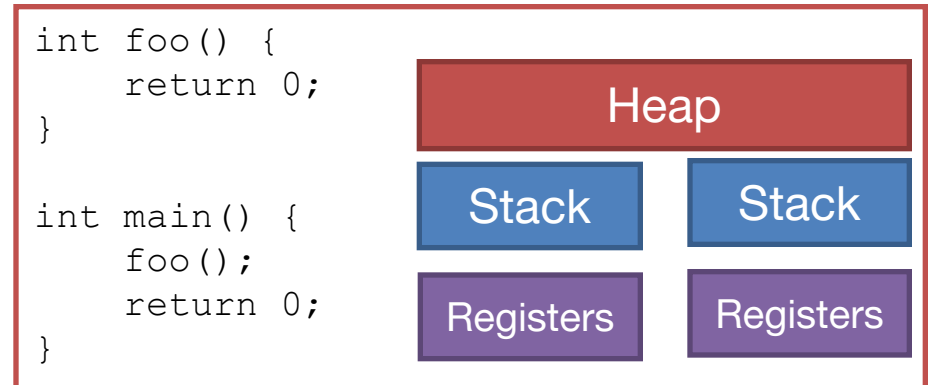
Threads and Processes

- Process: Execution stream + Program State
- Thread: Execution stream + Thread State (in same process address space)
 - “Lightweight process”

Process




Threads



- Can have multiple threads within a process

Processes and Threads



Thread State

Per process items

Address space
Page table
Open files
Child processes
Signals & handlers
Heap space
Global variables

Per thread items

Program counter
Registers
Stack & stack pointer
State

Per thread items

Program counter
Registers
Stack & stack pointer
State

Per thread items

Program counter
Registers
Stack & stack pointer
State

Thread Pros/ Cons

- Pros

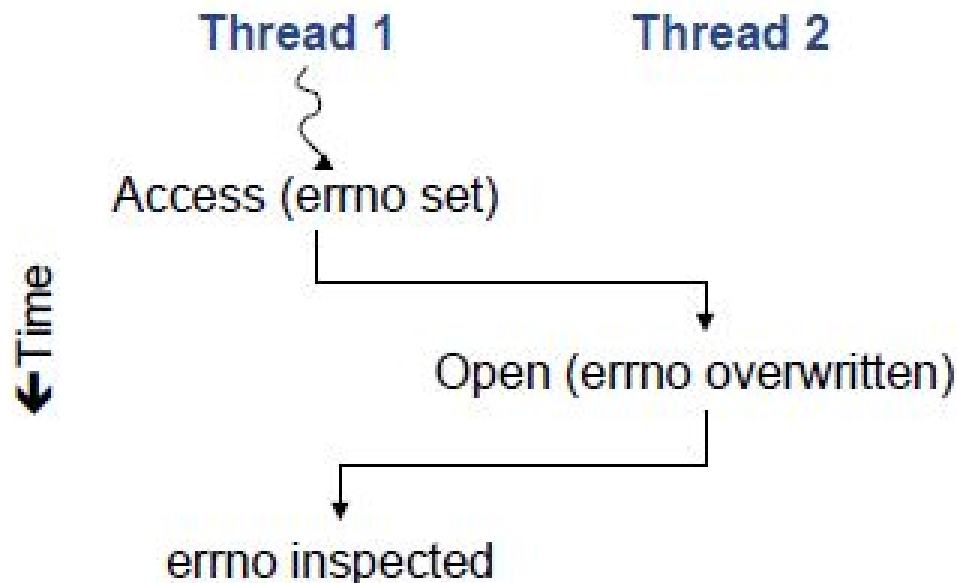
- Less memory to maintain Thread State (mostly extra stack space)
- Less performance overhead
- Automatic data sharing (same address space)

- Cons

- Shared memory can be source of data races
- Less robust compared to processes:
 - Memory corruption on one thread corrupts entire program
 - A crash in one thread crashes the entire program

Data Races in Multithreading

- Global variables and heap locations are problems
 - Not just your global variables, also library variables
- File operations modify **errno** to indicate errors
 - What if two threads both operate on separate files?



Re-entrant Code

- *Re-entrant*: ability to handle multiple calls to the same function at the same time
- Many libraries are not re-entrant
- Example a library that formats and prints messages to a log file:

```
int print_log(char * msg) {  
    fprintf(logfile, "%s", get_timestamp());  
    fprintf(logfile, msg);  
    fprintf(logfile, "\n");  
}
```


Thread Implementation

User-Level Threading vs.
Kernel-Level Threading

Who does thread management?

- User-level threading
 - User-level threading library
 - On thread create
 - Reserve new stack space using `mmap()`
 - On context-switch
 - Save / restore processor registers
 - Change stack pointer registers to stack of new thread
- Kernel-level threading
 - Kernel maintains thread state internally
 - Use system calls for thread creation / management

User Threads vs. Kernel Threads



User-level vs. Kernel-level

- User-level threading can be faster:
 - Do not need system calls for thread management
 - Context switching happens in user land
 - Can make application specific scheduling decisions for a tighter schedule
- However some issues with user-level:
 - Cannot run on multiple processors
 - To the kernel, it is still just one process
 - One thread can hog CPU with no OS intervention
 - If a thread performs I/O and blocks, all threads stop
 - The entire process transitions into a blocked state

CPU Hogging Solution

- Thread library provides an *yield()* function that allows thread to voluntarily give up CPU
 - Also called *cooperative threading*
 - Since within same process, no worry about malicious hogging by a third party
 - Voluntary yielding often results in a better schedule than blind OS scheduling
- To pre-empt thread, register a timer signal to fire after time slot
 - Perform context switch in signal handler

Blocking I/O Solution

- Avoid performing blocking I/O in threads
- Threading library provides own version of non-blocking I/O calls
 - Poll() and select() system calls inform whether an I/O on a file descriptor would block or not
 - Always check with poll() or select() before performing actual I/O
 - If I/O will block, sleep thread and schedule another thread

Multiple Processors Solution?

- None: Can't allocate CPU with no OS involvement
- Prevents user threading from being employed in an environment where parallelism is important (e.g. Servers)
- For small uniprocessor systems, it's perfectly fine (e.g. Embedded Systems)
 - Memory efficiency of user threading a plus
 - Fast context switches and cooperative scheduling help meet real-time constraints
 - Works even on simple kernels with no threading support

OS Threading Requirements

- User-level threading:
 - Support for non-blocking I/O in OS
- Kernel-level threading:
 - Thread management implementation in OS
- Most general purpose kernels support kernel threading
 - Linux Native POSIX Thread Library
 - Windows Threading API
- How about Java? Usually the Java Virtual Machine...
 - Maps threads to kernel-threads if it's supported
 - If no support for kernel-threads, uses user-level threading