

# Synchronization and Deadlocks

(or The Dangers of Threading)

CS449 Spring 2016

# Race Condition

Shared Data:

6
---

tail

1	8	5	6	20	?		
---	---	---	---	----	---	--	--

A[]

Enqueue():

thread switch

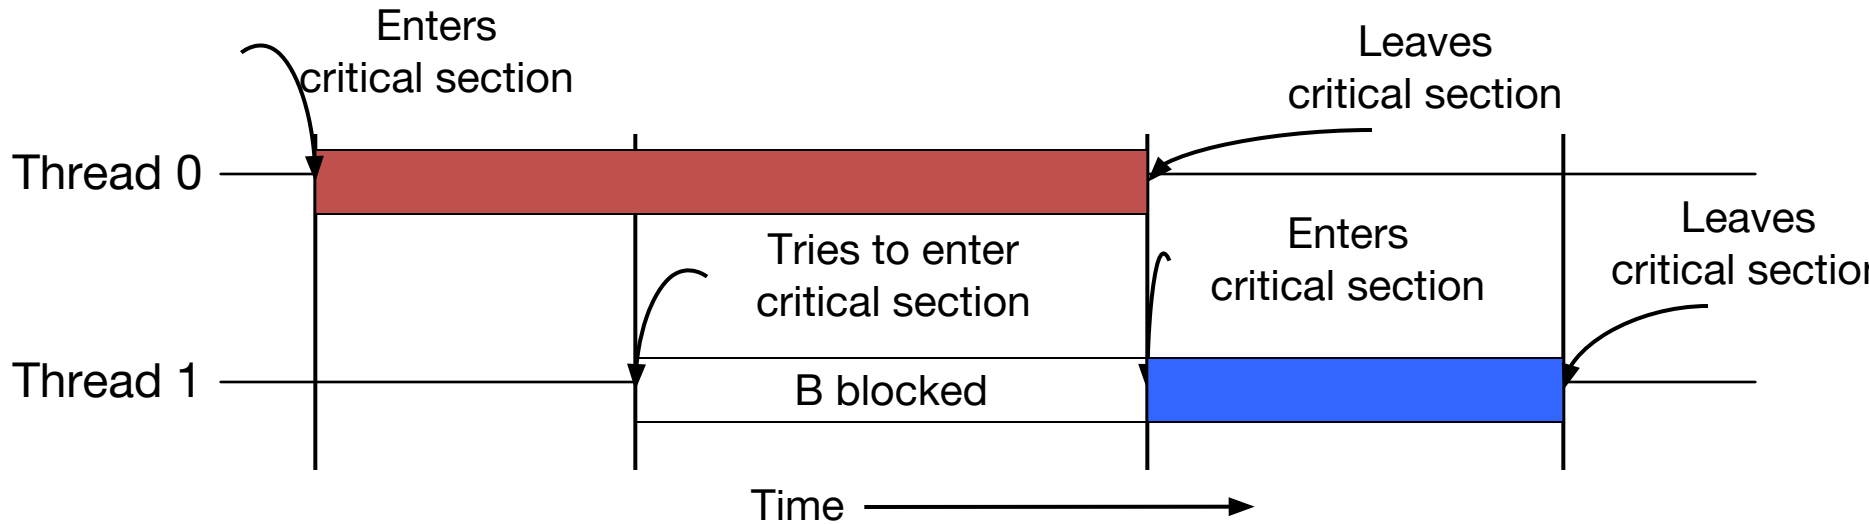
▶ A[tail] = 20;
▶ tail++;

Thread 0

▶ A[tail] = 9;
▶ tail++;

Thread 1

# Critical Sections



# Synchronization


- Scheduling can be random and preemption can happen at any time
- Two threads can be running in parallel
- Need a way to make critical sections “atomic”
  - Atom: smallest unit that cannot be divided further using chemical means
  - Atomic: section of code appears to execute as a unit with no interleaving
- Need help from the Operating System (or the user thread scheduler)

# Mutex


- MUTual EXclusion
- A mutex is a lock that only one thread can acquire
- All other threads attempting to enter the critical section will be blocked
- Thread holding lock releases lock when exiting critical section

# Critical Sections

Shared Data:

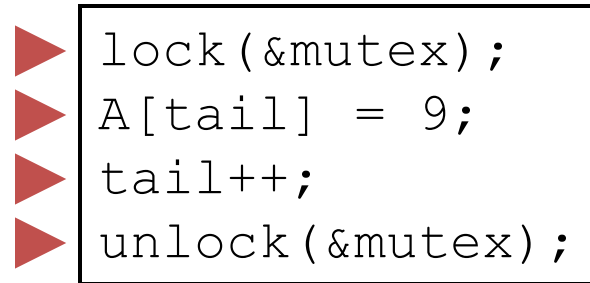


Enqueue():



Thread 0

**Blocked!**



Thread 1

# pthread\_mutex\_t


```
#include <stdio.h>
#include <pthread.h>

int tail = 0;
int A[20];

pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;

void enqueue(int value)
{
    pthread_mutex_lock(&mutex);
    A[tail] = value;
    tail++;
    pthread_mutex_unlock(&mutex);
}
```

# Sharing can lead to **deadlocks**!





# Deadlocks


- “A set of processes is **deadlocked** if each process in the set is waiting for an event that only another process in the set can cause.”

# Deadlock Requirements

- Mutual Exclusion
  - Resource can only be held by one process at a time
- Hold and Wait
  - Process does not release its resources when waiting for other resources
- No Preemption
  - Resources cannot be forcibly taken away
- Circular Wait
  - Processes are waiting for each other in a closed cycle


# Deadlock Example: Dining Philosophers

- Philosophers eat/think
- Eating needs 2 forks
- Pick one fork at a time
- What happens if each philosopher grabs the fork on the right?
  1. Mutual exclusion
  2. Hold and wait
  3. No preemption of resource
  4. *Circular wait*



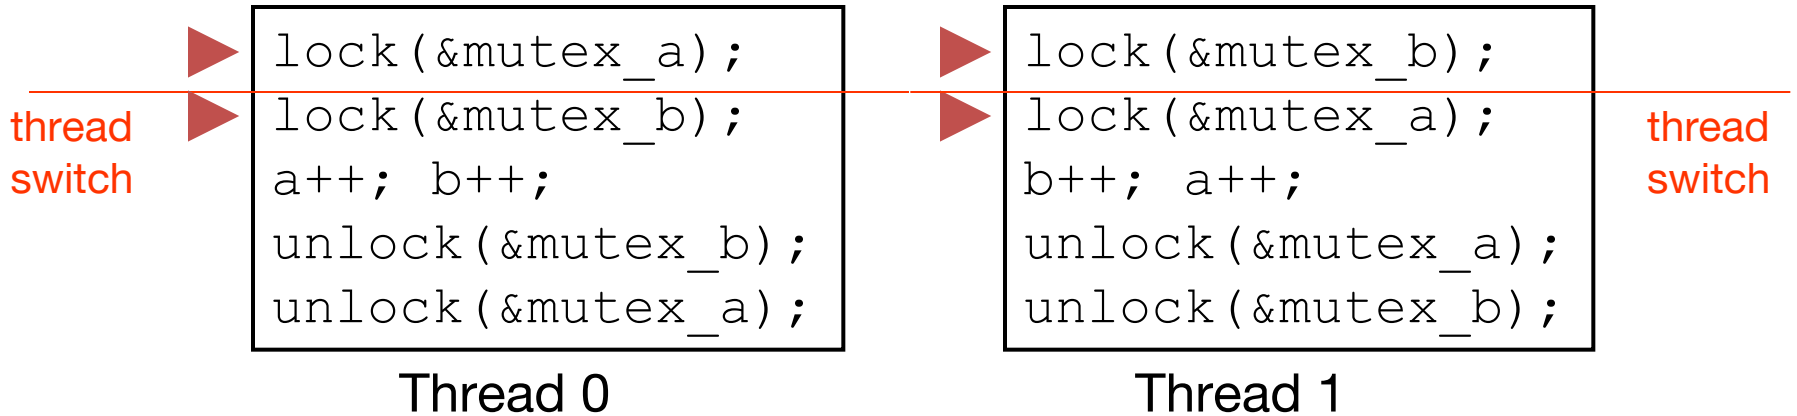
# Deadlock Example

Shared Data:



**Blocked!**

**Blocked!**



# Handling Deadlocks

Detect


Prevent

Avoid

Ignore

# Deadlock Solution 1: Remove Circular Wait

Shared Data:



```
lock(&mutex_a);  
lock(&mutex_b);  
a++; b++;  
unlock(&mutex_b);  
unlock(&mutex_a);
```

Thread 0


```
lock(&mutex_a);  
lock(&mutex_b);  
b++; a++;  
unlock(&mutex_b);  
unlock(&mutex_a);
```

Thread 1

- Rule: Acquire locks in the same order

# Deadlock Solution 2: Remove Hold and Wait

Shared Data:



```
lock(&mutex_a);  
a++;  
unlock(&mutex_a);  
lock(&mutex_b);  
b++;  
unlock(&mutex_b);
```

Thread 0

```
lock(&mutex_b);  
b++;  
unlock(&mutex_b);  
lock(&mutex_a);  
a++;  
unlock(&mutex_a);
```

Thread 1

# Producer/Consumer Problem

## Shared variables

```
#define N 10;

int buffer[N];
int in = 0, out = 0, counter = 0;
```

## Producer

```
while (1) {
    if (counter == N)
        sleep();

    buffer[in] = ... ;
    in = (in+1) % N;

    counter++;

    if (counter==1)
        wakeup(consumer);
}
```

## Consumer

```
while (1) {
    if (counter == 0)
        sleep();

    ... = buffer[out];
    out = (out+1) % N;

    counter--;

    if (counter == N-1)
        wakeup(producer);
}
```



# Condition Variables

- A condition under which a thread executes or is blocked
- `pthread_cond_t`
  - The type of the condition variable
- `pthread_cond_wait (condition, mutex)`
  - Blocks current thread until condition signaled
- `pthread_cond_signal (condition)`
  - Unblocks one thread waiting for condition

# Producer/Consumer

```
#define N 10
int buffer[N];
int counter = 0, in = 0, out = 0, total = 0;

pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t prod_cond = PTHREAD_COND_INITIALIZER;
pthread_cond_t cons_cond = PTHREAD_COND_INITIALIZER;

void *producer(void *junk) {
    while(1) {
        pthread_mutex_lock(&mutex);
        if( counter == N )
            pthread_cond_wait(&prod_cond,
                              &mutex);

        buffer[in] = total++;
        printf("Produced: %d\n",
               buffer[in]);

        in = (in + 1) % N;
        counter++;

        if( counter == 1 )
            pthread_cond_signal(&cons_cond);

        pthread_mutex_unlock(&mutex);
    }
}

void *consumer(void *junk) {
    while(1) {
        pthread_mutex_lock(&mutex);
        if( counter == 0 )
            pthread_cond_wait(&cons_cond,
                              &mutex);

        printf("Consumed: %d\n",
               buffer[out]);

        out = (out + 1) % N;
        counter--;

        if( counter == (N-1) )
            pthread_cond_signal(&prod_cond);

        pthread_mutex_unlock(&mutex);
    }
}
```

# Semaphores

- Mutex: controls access to one resource
  - Critical section protects access to a shared data structure
  - Analogy: Room with one desk where only one student can enter
- Semaphore: controls access to multiple resources
  - Analogy: Room with 5 desks where up to 5 students can enter
  - Internal counter keeps track of number of resources available
- Mutexes can be thought of as binary semaphores (semaphores that can only count up to 1 resource)

# Producer/Consumer

```
#include <semaphore.h>
```

```
#define N 10
```

```
int buffer[N];
```

```
int in = 0, out = 0, total = 0;
```

```
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
```

```
sem_t semfull; // sem_init(&semfull, 0, 0); in main()
```

```
sem_t semempty; // sem_init(&semempty, 0, N); in main()
```

```
void *producer(void *junk) {  
    while(1) {  
        sem_wait(&semempty);  
        pthread_mutex_lock(&mutex);  
  
        buffer[in] = total++;  
        printf("Produced: %d\n",  
               buffer[in]);  
  
        in = (in + 1) % N;  
  
        pthread_mutex_unlock(&mutex);  
        sem_post(&semfull);  
    }  
}
```

```
void *consumer(void *junk) {  
    while(1) {  
        sem_wait(&semfull);  
        pthread_mutex_lock(&mutex);  
  
        printf("Consumed: %d\n",  
               buffer[out]);  
  
        out = (out + 1) % N;  
  
        pthread_mutex_unlock(&mutex);  
        sem_post(&semempty);  
    }  
}
```

# Producer/Consumer

```
#include <semaphore.h>
```

```
#define N 10
```

```
int buffer[N];
```

```
int in = 0, out = 0, total = 0;
```

```
sem_t semmutex; // sem_init(&semmutex, 0, 1); in main()
```

```
sem_t semfull; // sem_init(&semfull, 0, 0); in main()
```

```
sem_t semempty; // sem_init(&semempty, 0, N); in main()
```

```
void *producer(void *junk) {  
    while(1) {  
        sem_wait(&semempty);  
        sem_wait(&semmutex);  
  
        buffer[in] = total++;  
        printf("Produced: %d\n",  
               buffer  
[in]);  
        in = (in + 1) % N;  
  
        sem_post(&semmutex);  
        sem_post(&semfull);  
    }  
}
```

```
void *consumer(void *junk) {  
    while(1) {  
        sem_wait(&semfull);  
        sem_wait(&semmutex);  
  
        printf("Consumed: %d\n",  
               buffer[out]);  
        out = (out + 1) % N;  
  
        sem_post(&semmutex);  
        sem_post(&semempty);  
    }  
}
```

# Deadlock!

```
#include <semaphore.h>
```

```
#define N 10
```

```
int buffer[N];
```

```
int in = 0, out = 0, total = 0;
```

```
sem_t semmutex; // sem_init(&semmutex, 0, 1); in main()
```

```
sem_t semfull; // sem_init(&semfull, 0, 0); in main()
```

```
sem_t semempty; // sem_init(&semempty, 0, N); in main()
```

```
void *producer(void *junk) {  
    while(1) {  
        sem_wait(&semmutex);  
        sem_wait(&semempty);  
  
        buffer[in] = total++;  
        printf("Produced: %d\n",  
               buffer  
[in]);  
        in = (in + 1) % N;  
  
        sem_post(&semfull);  
        sem_post(&semmutex);  
    }  
}
```

```
void *consumer(void *junk) {  
    while(1) {  
        sem_wait(&semmutex);  
        sem_wait(&semfull);  
  
        printf("Consumed: %d\n",  
               buffer[out]);  
        out = (out + 1) % N;  
  
        sem_post(&semempty);  
        sem_post(&semmutex);  
    }  
}
```

# Valgrind

- Same tool we used for memory errors
- Helgrind: component of valgrind that does potential data race / deadlock detection
- Command: `valgrind --tool=helgrind <program>`
- Not perfect.
  - Can miss errors (sometimes)
  - Can report errors when there are none (sometimes)
- Not a replacement for sound programming