


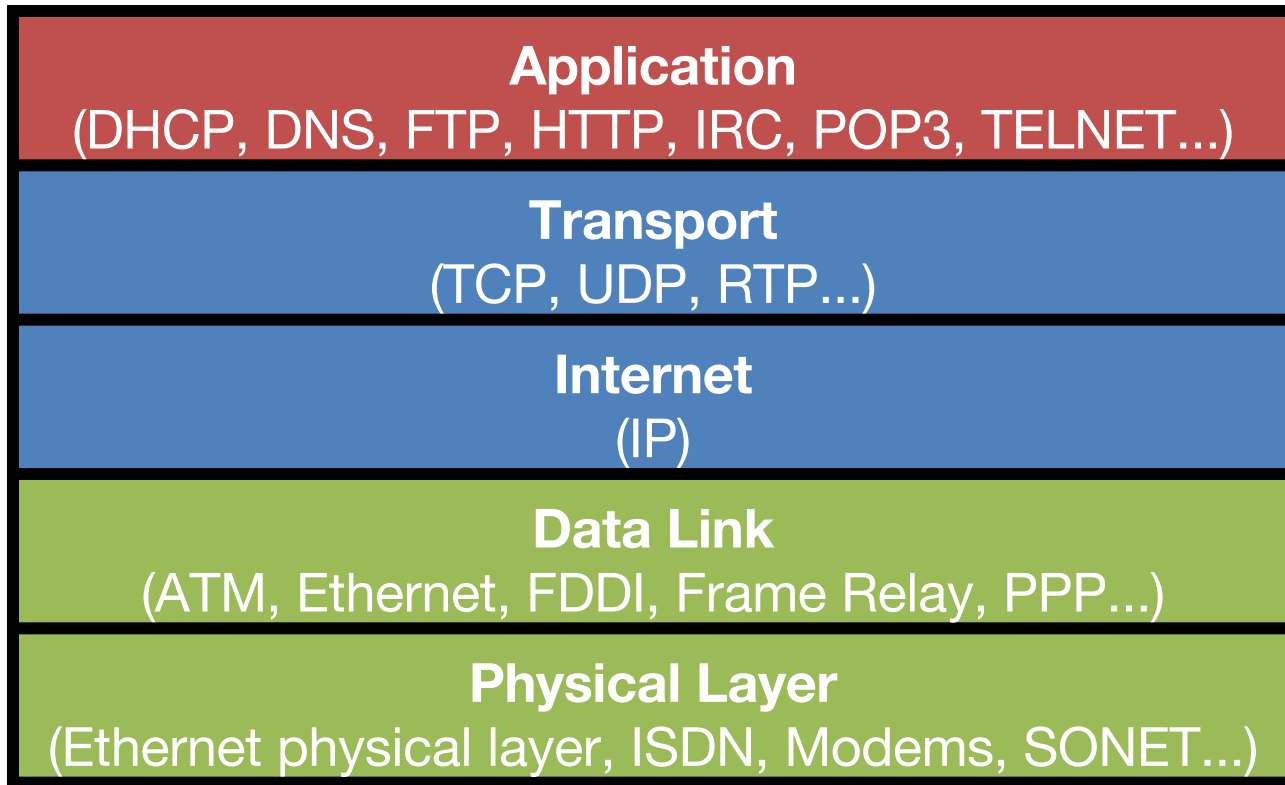
Network Communication

CS449 Spring 2016

Network



Internet Layer Model



Internet Protocol (IP)

- **Protocol** – A standard procedure for regulating data transmission between computers.
- **IP Addresses** – 32-bit (v4) or 128 (v6) number denoting an destination or source
- **Port** – a number representing a particular listener on a machine

Transmission Control Protocol (TCP)

- **Connection-oriented** – Make a circuit with a remote machine
- Guarantees data arrives, and in-order


User Datagram Protocol (UDP)

- **Datagram** – (play on telegram) A message with no acknowledgement
- **Connectionless** – Send and forget
- No guarantee data arrives or is in the order sent

IP Packet

| | Bits 0–3 | 4–7 | 8–15 | 16–18 | 19–31 |
|------|---------------------|---------------|-----------------|-----------------|-----------------|
| 0 | Version | Header length | Type of Service | Total Length | |
| 32 | Identification | | | Flags | Fragment Offset |
| 64 | Time to Live | | Protocol | Header Checksum | |
| 96 | Source Address | | | | |
| 128 | Destination Address | | | | |
| 160 | Options | | | | |
| 192- | Data | | | | |

Packets



NETWORK PROGRAMMING

Processes communicating

Process: program running within a host.

- within same host, two processes communicate using **inter-process communication** (defined by OS).
- processes in different hosts communicate by exchanging **messages**

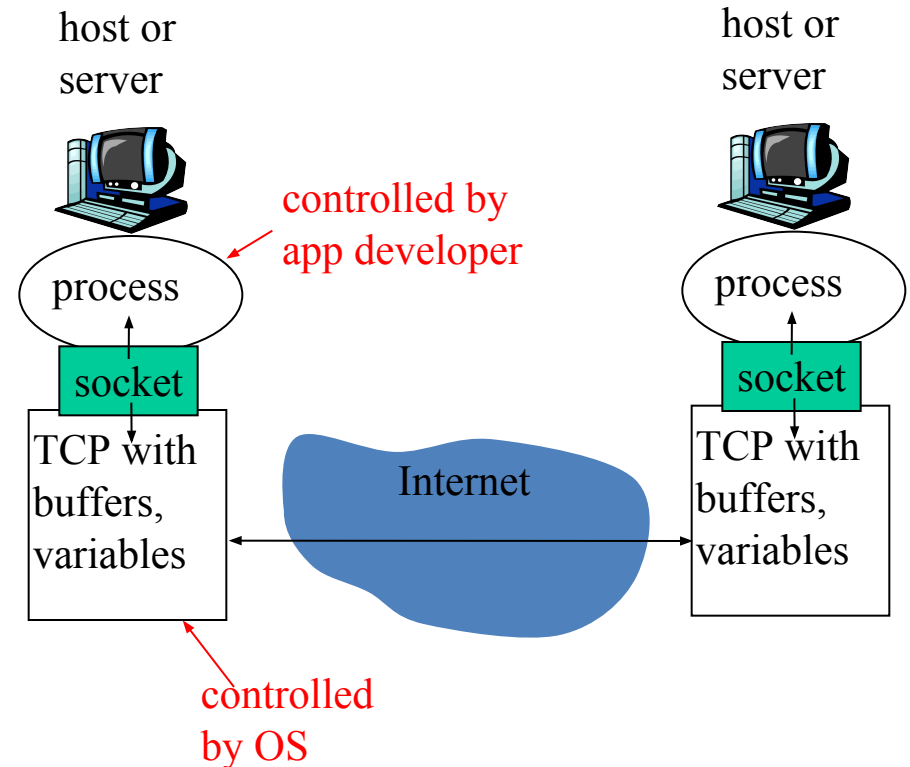
Client process: process that initiates communication

Server process: process that waits to be contacted

- Note: applications with P2P architectures have client processes & server processes

Sockets

- ❑ process sends/receives messages to/from a **socket**
- ❑ socket is a software communication channel
 - ❑ sending process sends into socket
 - ❑ sending process relies on transport infrastructure on other side of socket to deliver message to socket at receiving process



API: (1) choice of transport protocol; (2) ability to fix a few parameters (lots more on this later)

Addressing processes

- ❑ For a process to receive messages, it must have an identifier
- ❑ A host has a unique 32-bit IP address

Q: does the IP address of the host on which the process runs suffice for identifying the process?

Answer: No, many processes can be running on same host

- ❑ Identifier includes both the IP address and **port numbers** associated with the process on the host.
- ❑ Example port numbers:
 - ❑ HTTP server: 80
 - ❑ Mail server: 25

App-layer protocol defines

- ❑ Types of messages exchanged,
 - ❑ eg, request & response
- ❑ Message Syntax
 - ❑ What fields in messages & how fields are **delineated**
- ❑ Message Semantics
 - ❑ Meaning of information in fields
- ❑ Rules for when and how processes send & respond to messages

Public-domain protocols:

- ❑ defined in RFCs
- ❑ allows for interoperability
- ❑ eg, HTTP, SMTP

Proprietary protocols:

- ❑ eg, Skype

BERKELEY SOCKETS

Socket

- UNIX treats everything as a file
 - File Descriptor
 - read()/ write()
- Treat network as a file called a socket
- Berkeley sockets are de facto standard API

socket()

- Creates a Socket Descriptor

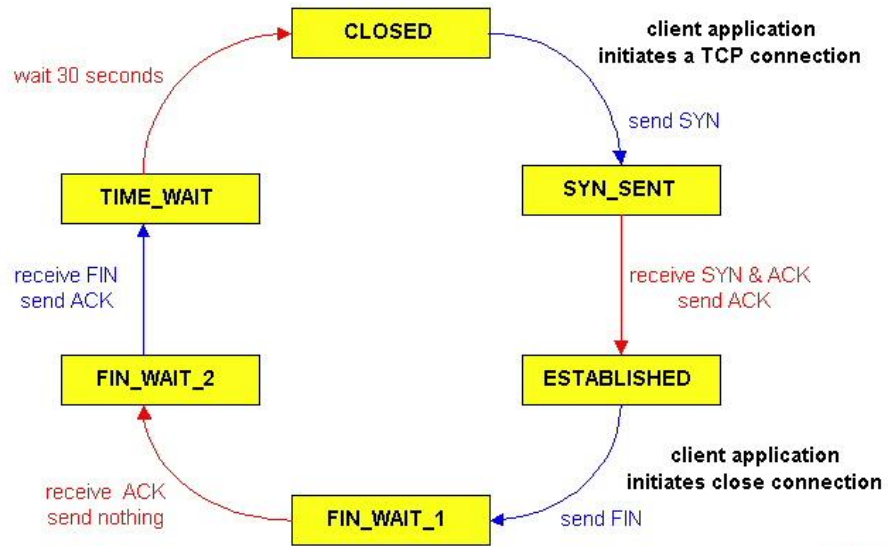
```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol);
```

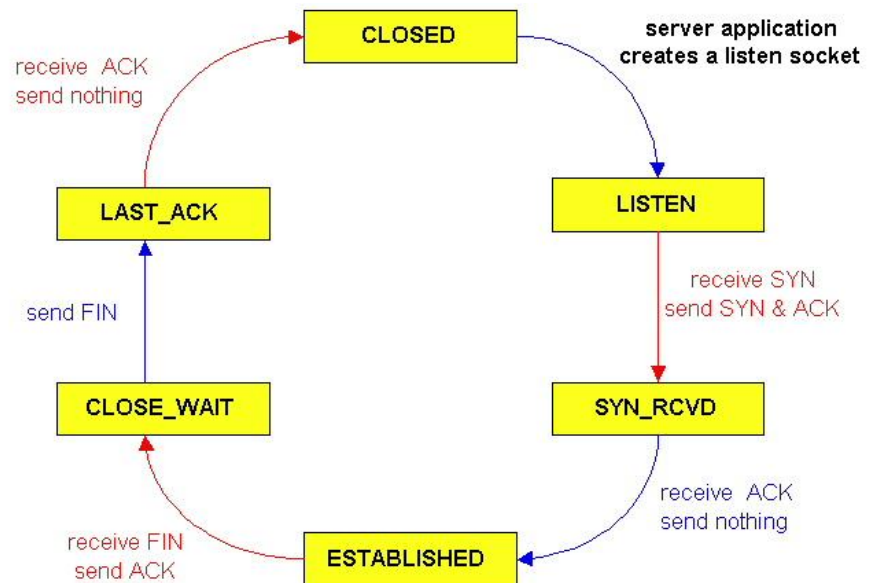
| Parameter | Values |
|-----------|---|
| domain | •AF_INET for IPv4 •AF_INET6 for IPv6 |
| type | •SOCK_STREAM •SOCK_DGRAM •SOCK_SEQPACKET •SOCK_RAW |
| protocol | IPPROTO_IP (defined as 0) |

TCP Connection Management



TCP client lifecycle

TCP server lifecycle



SERVER STUFF

bind()

- **Attach a socket to a port**

```
int bind(int sockfd, struct sockaddr *addr, int  
        addrlen);
```

```
struct sockaddr_in addr;  
memset(&addr, 0, sizeof(addr));  
addr.sin_family = AF_INET;  
addr.sin_port = htons(PORT);  
addr.sin_addr.s_addr = inet_addr("127.0.0.1");  
bind(fd, (struct sockaddr *)addr, sizeof(addr));
```

- **Returns 0 on success, -1 on error**

listen()

- Set up a listening socket

```
int listen(int sockfd, int backlog);
```

- Backlog – how many pending connections are allowed to wait (OS max usually around 20, set to lower, ~10)
- Returns 0 on success, -1 on error

accept()

- Block and wait for connection to occur

```
int accept(int sockfd, struct sockaddr  
          *cliaddr, socklen_t *addrlen);
```

- Returns file descriptor for the new socket, or -1 on error
- Will return information about the client through the structure (can be NULL)

send() and recv()

```
int send(int sockfd, const void *msg, int  
len, int flags);
```

- By default blocks until `len` bytes are sent
- Returns number of bytes sent, or -1 on error

```
int recv(int sockfd, void *buf, int len,  
unsigned int flags);
```

- Returns any amount of data available up to `len` bytes
- Returns number of bytes received, or 0 if peer closed connection, or -1 on error

errno and perror()

- Whenever operations on descriptors fail, the `errno` global variable is set with error type
- Especially important to check for error for socket operations because it involves remote machine

- `void perror(const char *s)`

- prints descriptive message to stdout for `errno`, prefixed by `s` (if not NULL)

```
if (bind(socket, &addr, sizeof(addr)) < 0
{
    perror(NULL) ;
}
```

CLIENT STUFF

connect()

- Connect to a server located at some address and port

```
int connect(int sockfd, struct sockaddr *serv_addr,  
            int addrlen);
```

```
struct sockaddr_in addr;  
memset(&addr, 0, sizeof(addr));  
addr.sin_family = AF_INET;  
addr.sin_port = htons(PORT);  
addr.sin_addr.s_addr = net_addr("127.0.0.1");  
connect(fd, (struct sockaddr *)&addr, sizeof(addr));
```

send() and recv()

```
int send(int sockfd, const void *msg, int  
len, int flags);
```

```
int recv(int sockfd, void *buf, int len,  
unsigned int flags);
```

CONNECTIONLESS COMMUNICATION

Datagram Send and Receive

```
int sendto(int sockfd, const void *msg, int  
    len, unsigned int flags, const struct  
    sockaddr *to, socklen_t tolen);
```

```
int recvfrom(int sockfd, void *buf, int len,  
    unsigned int flags, struct sockaddr *from,  
    int *fromlen);
```

DNS

- Domain Name Server
- Resolve a name to an IP address:

<http://www.cs.pitt.edu> -> 130.49.220.23

DNS

```
#include <netdb.h>

struct hostent *gethostbyname(const char *name);

struct hostent {
    char    *h_name;           /* official name of host */
    char    **h_aliases;       /* alias list */
    int     h_addrtype;        /* host address type */
    int     h_length;          /* length of address */
    char    **h_addr_list;     /* list of addresses */
}

#define h_addr h_addr_list[0] /* for backward
    compatibility */
```